

High Performance & Scalable Analytics

HPSA – 2019

Master in Big Data Analytics & Social Mining

nicola.tonellotto@isti.cnr.it

Learning Objectives

- **Knowledge**
 - High performance computing
 - Systems & storage for Big Data
 - Programming for Big Data
- **Skills**
 - Interacting with distributed resources
 - Use of distributed storage systems
 - Use of distributed data processing tools
 - How to prepare big data for HPSA
 - How to process big data with HPSA
 - How to make inference from big data through HPSA
 - How to assess results from HPSA

Tentative Timetable

- **Wednesday 17/4/19 (9-13)**
 - Course presentation
 - Project discussion
 - Parallel/distributed computing
 - Map/Reduce
 - Remote access to resources
- **Friday 19/4/19 (9-13)**
 - Distributed filesystem
 - HDFS management
 - Data analytics with Spark I
- **Thursday 2/5/19 (9-13)**
 - Data analytics with Spark II
 - Machine learning with Spark I
- **Saturday 4/5/19 (9-13)**
 - Data analytics with Spark III
 - Machine learning with Spark II
- **Thursday 9/5/19 (9-12)**
 - Use case walkthrough
- **Thursday 16/5/19 (14-15)**
 - Final exam

Admin I

- Web site hosted on Github

<https://github.com/tonellotto/hpsa>

- Select the current year branch

<https://github.com/tonellotto/hpsa/tree/aa1819>

The screenshot shows the GitHub repository page for 'tonellotto / hpsa'. At the top, there are buttons for 'Unwatch' (1), 'Star' (0), and 'Fork' (0). Below the header, there are tabs for 'Code', 'Issues (0)', 'Pull requests (0)', 'Projects (0)', 'Wiki', 'Insights', and 'Settings'. The 'Code' tab is selected. In the 'Description' section, it says 'Course Material Repository for the High Performance & Scalable Analytics course of t' and has a 'Save' button. The 'Website' section has a placeholder 'Website for this repository (optional)' and a 'Save' button. Under 'Manage topics', there are sections for '2 commits', '2 branches', '0 releases', and '1 contributor'. A yellow box highlights the 'aa1819 (1 minute ago)' branch under 'Your recently pushed branches:' and a green 'Compare & pull request' button. On the left, a sidebar shows a dropdown for 'Branch: master' and a 'New pull request' button. Below that is a 'Switch branches/tags' section with a search bar 'Find or create a branch...', a 'Branches' tab, and a list of branches including 'aa1819' and 'master'. The 'master' branch is checked. The main right area shows a commit history with 'Latest commit b7eb0d0 16 minutes ago' and a file named 'Update README.md' updated '16 minutes ago'. At the bottom, there are buttons for 'Create new file', 'Upload files', 'Find File', and 'Clone or download'.

Admin II

```
$> # sudo apt-get install git  
$> cd $HOME  
$> git clone https://github.com/tonellotto/hpsa  
$> git checkout aa1819  
$> ls -ltrh
```

Admin III

- **Teacher:** Nicola Tonellotto (the bad cop)
- **Tutor:** Roberto Trani (the good cop)
- **Tools:** SSH/SCP, HDFS, Spark, MLLib
- **Credits:** 2
- **Hours:** 20



Projects

- All projects are composed by 4 phases:
 1. Data preparation
 2. Feature extraction
 3. Machine learning
 4. Results assessment
- Project topics
 - Spam classification of SMS
 - Sentiment analysis of movie reviews
 - Category identification of beers
 - Stock prediction using news
 - Customer segmentation of bank accounts
 - Customer segmentation from e-commerce transactions

Final Grading

- **50% project evaluation**

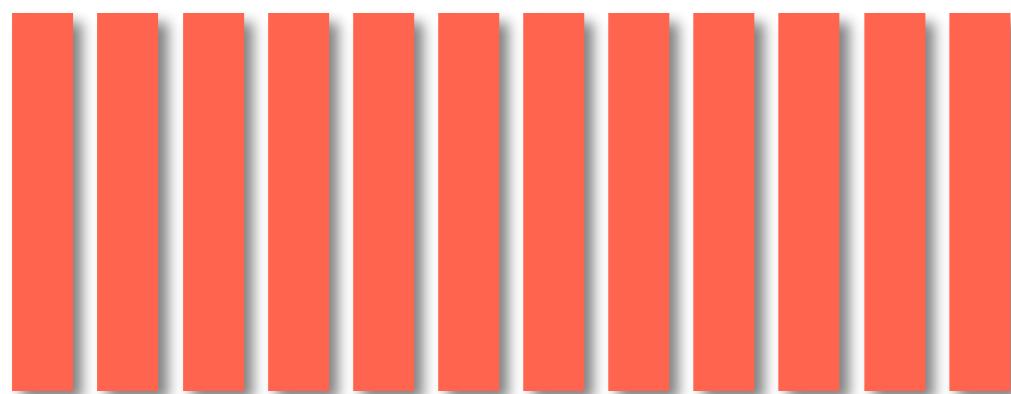
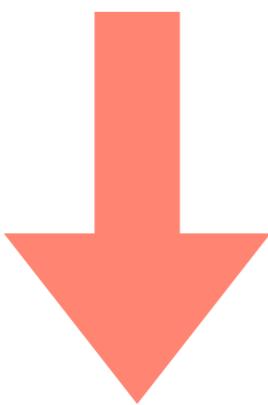
- Carried out during the open lab hours
- 2 pages report (font 11 pt, single interline, A4 standard margins)
- 1 self-contained python notebook
- Group activity (max 6 per project)
- Final version to be submitted by email to the tutor no later than 13/5/19, 11:59 AM CET
- *Group grade*

- **50% written exam**

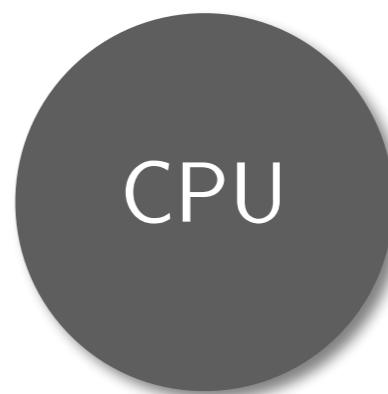
- Thursday 16/5/19, 13:45 in the classroom
- Multiple choice questions, open-ended questions, programming questions
- *Individual grade*

Classical Computation

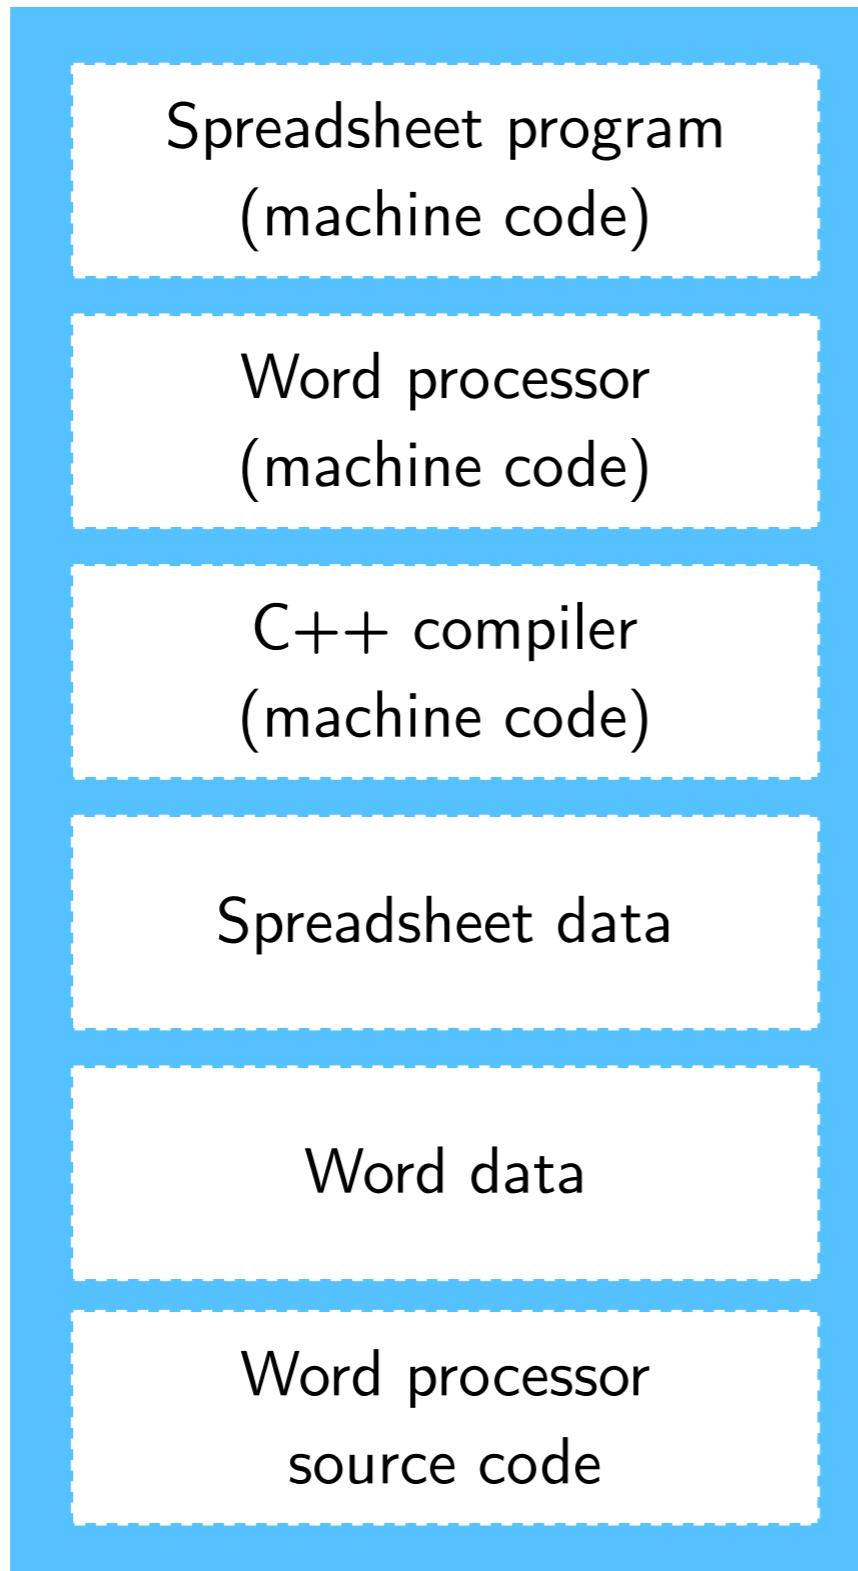
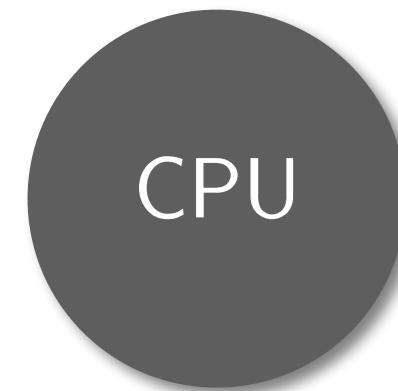
Algorithm



Instructions



Stored Program Computer

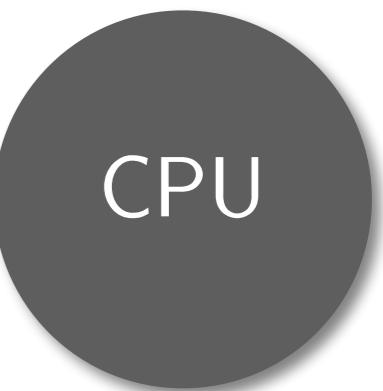
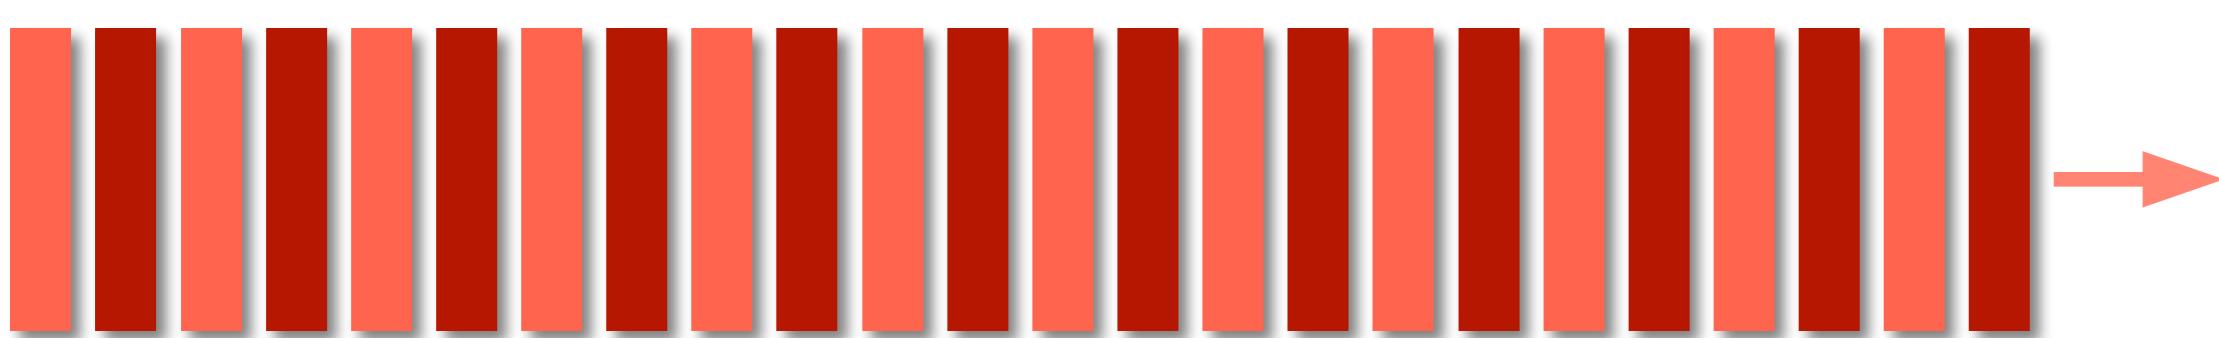
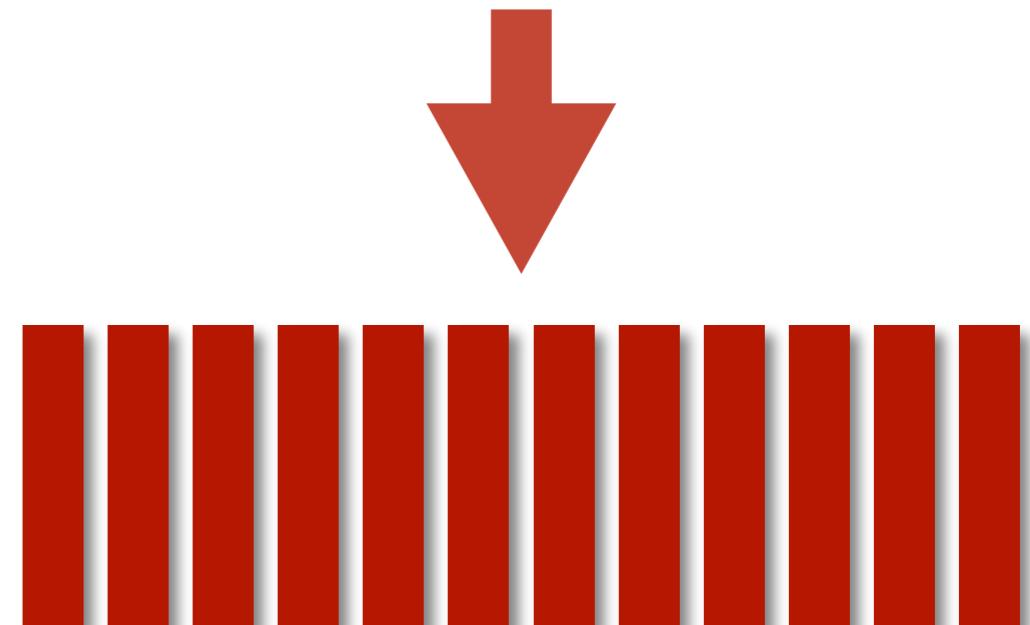
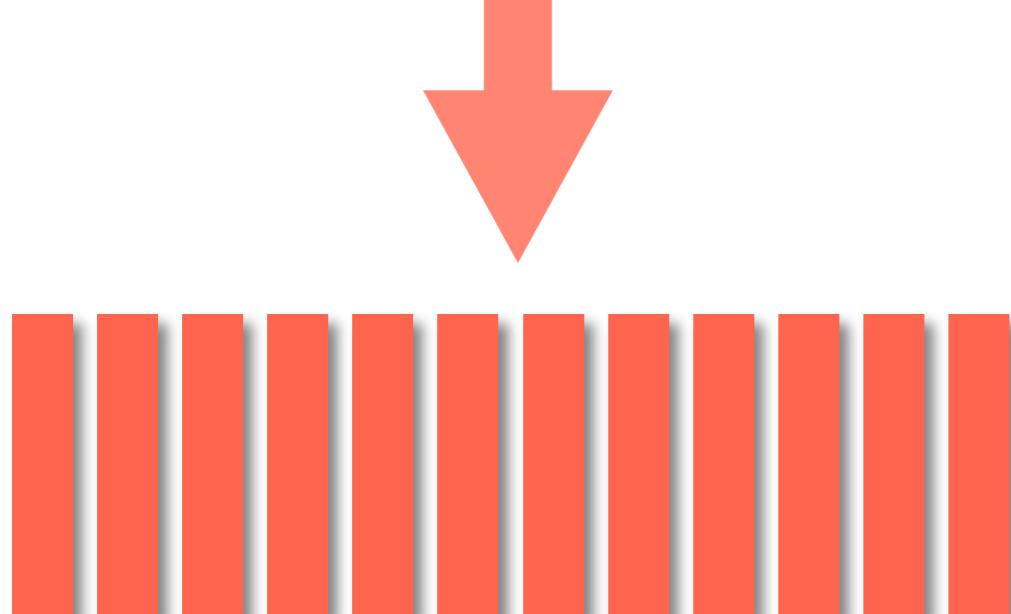


- Instructions represented as binary numbers (like data) (0010010101111001...)
- Data and instructions stored in main memory
- A CPU:
 - **FETCH**: get instructions from memory
 - **DECODE**: identify what to do
 - **EXECUTE**: perform actions

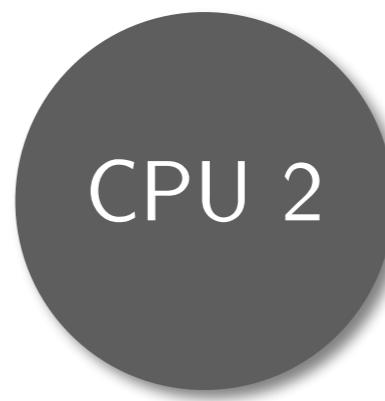
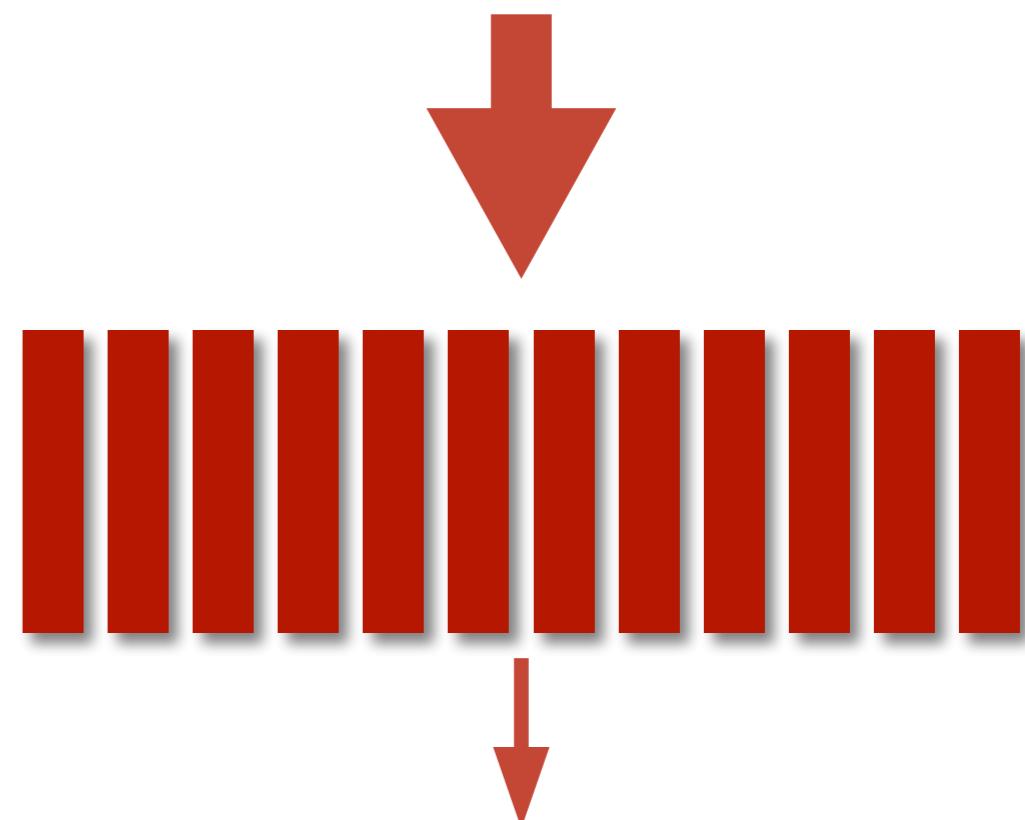
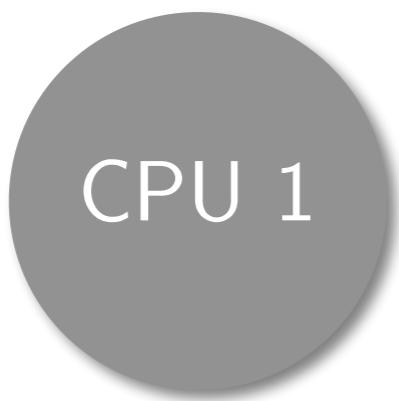
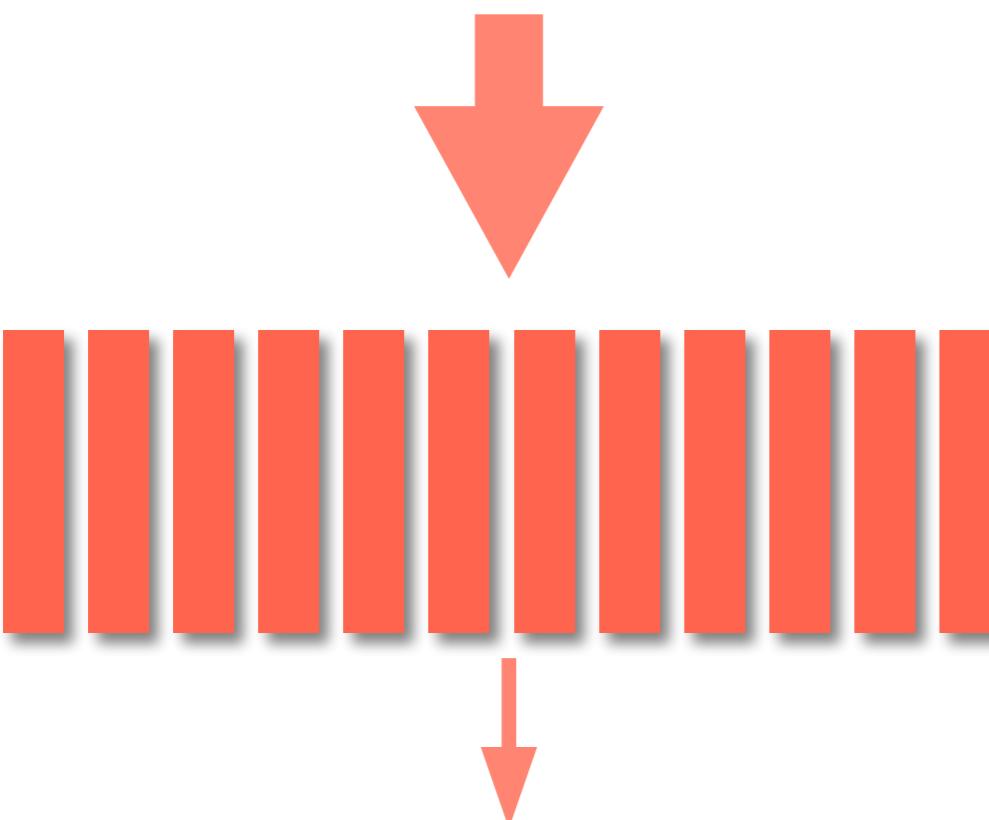
Time Sharing

Word processor

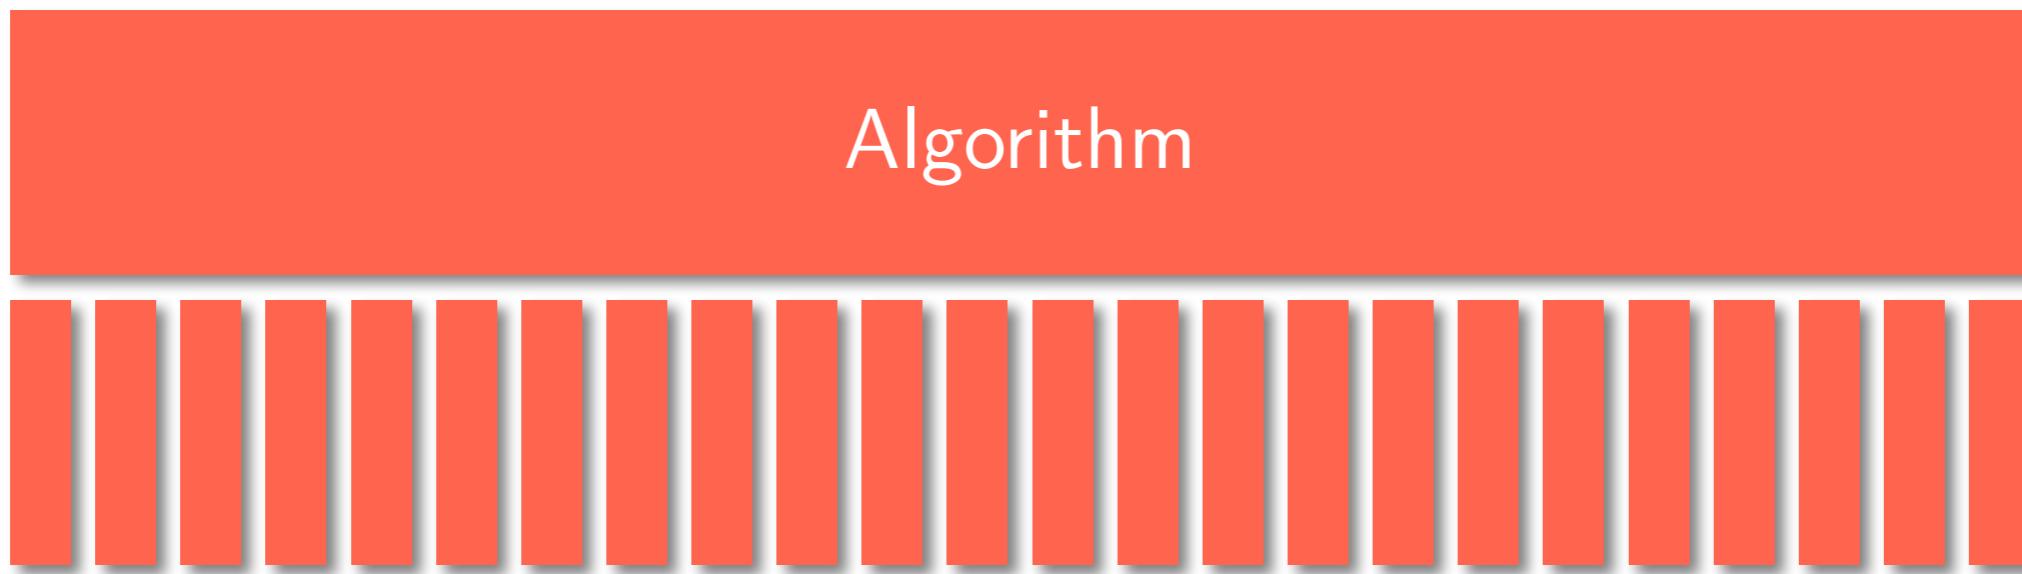
Spreadsheet



Parallel Computation

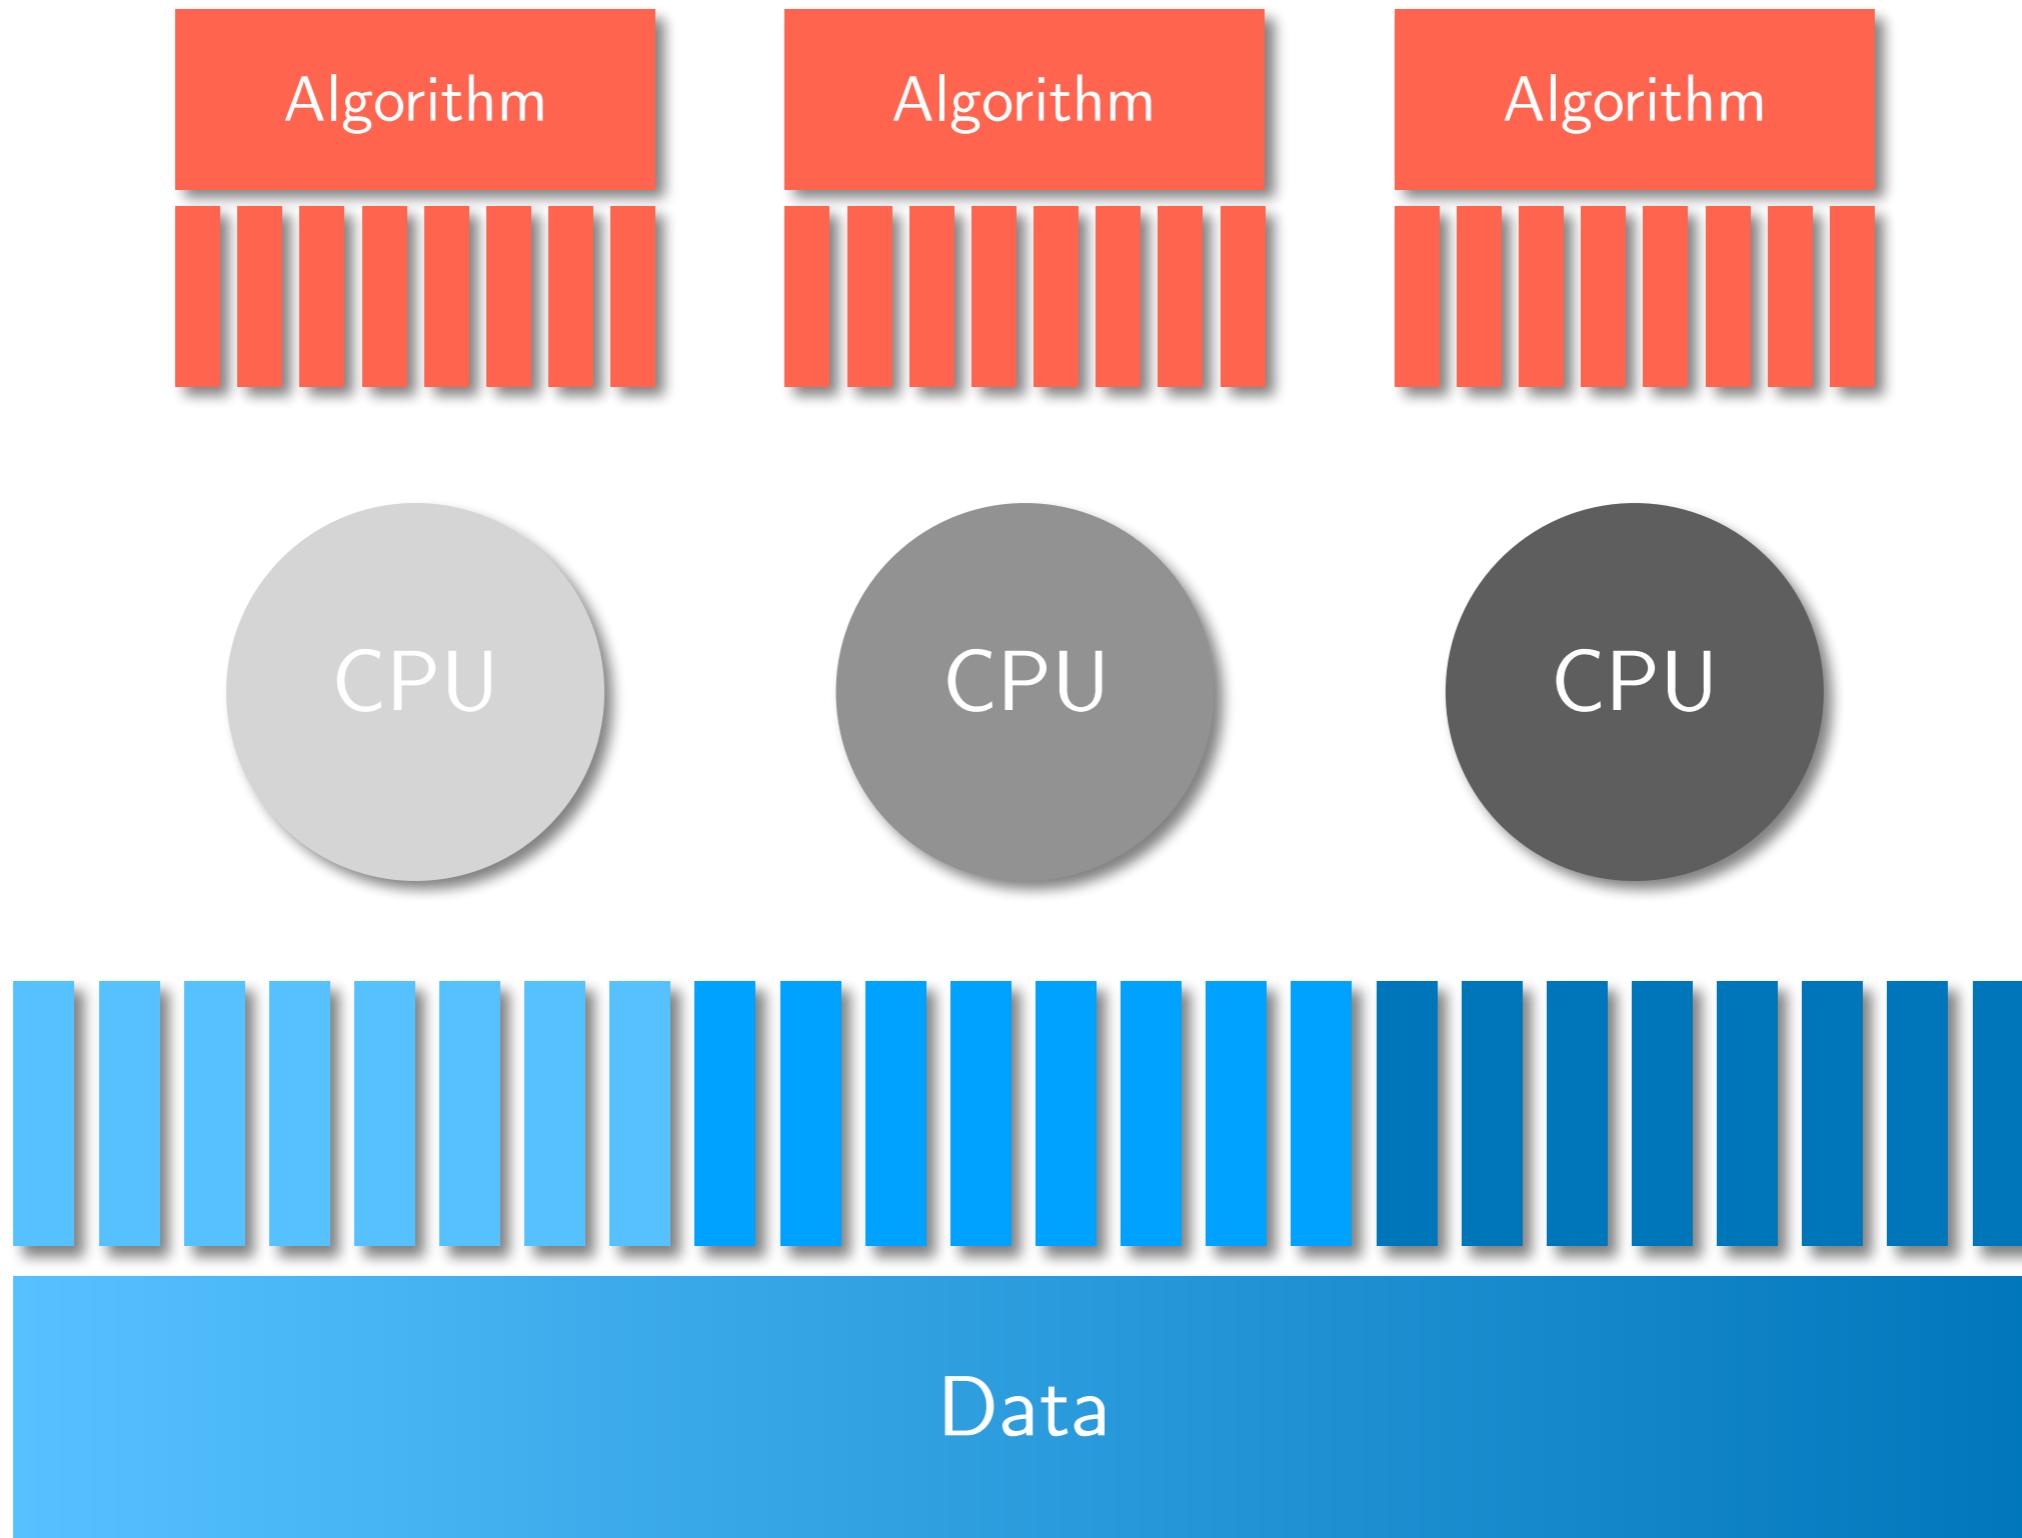


Single Instruction Single Data

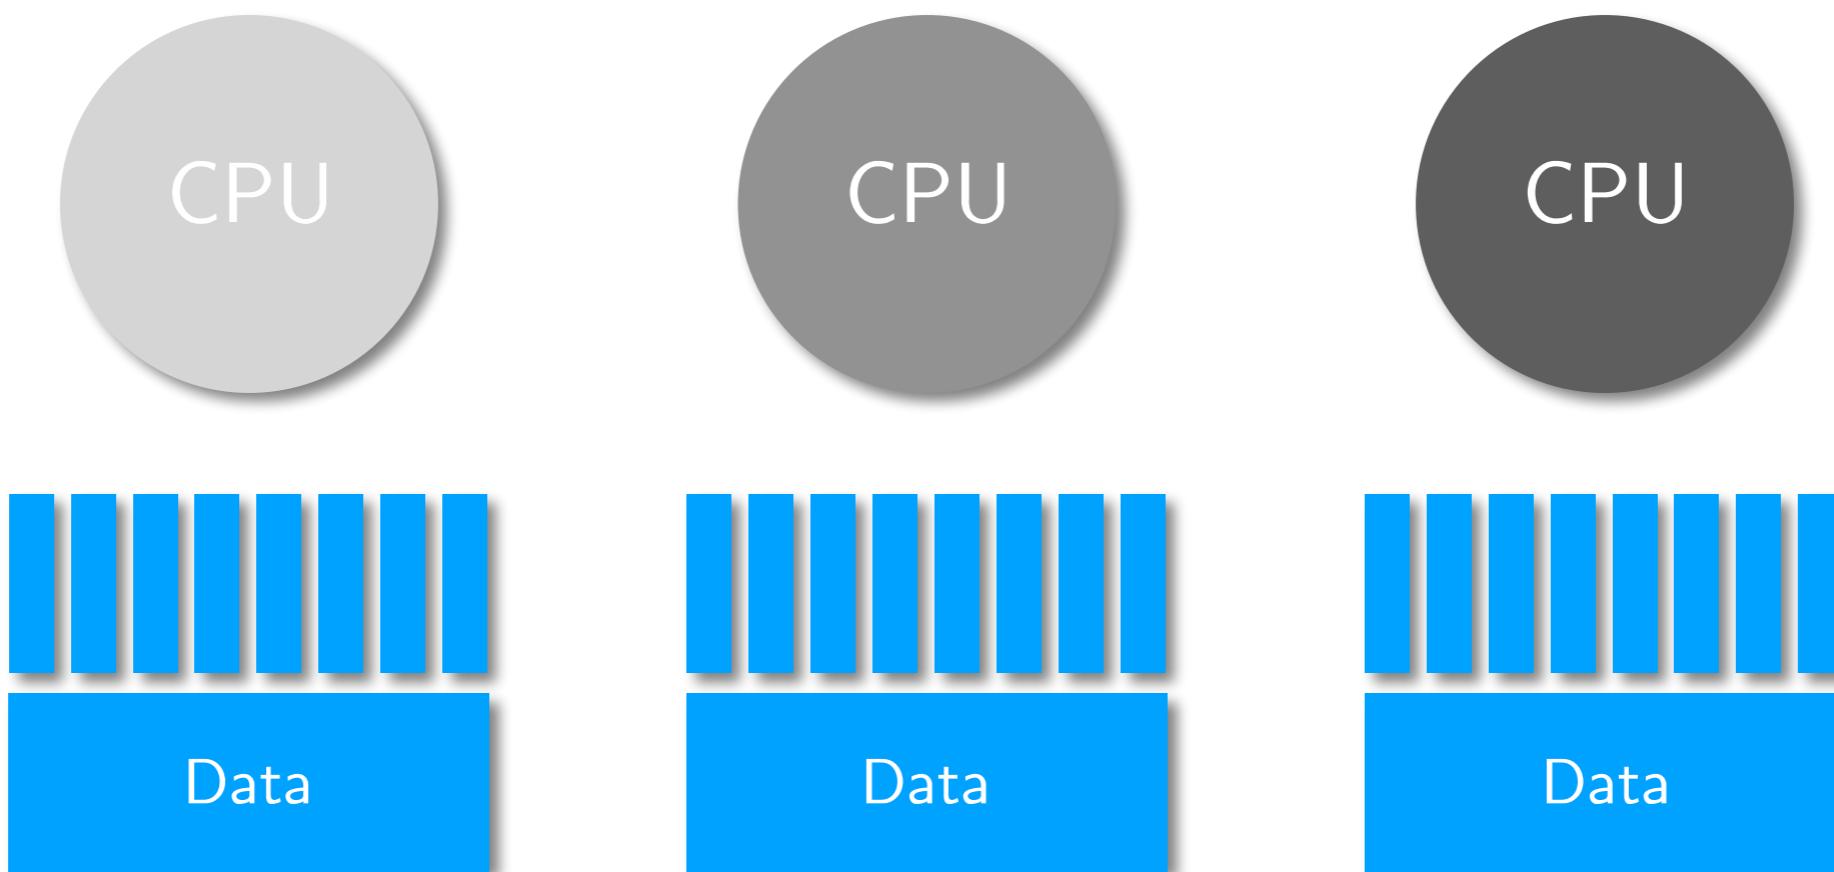
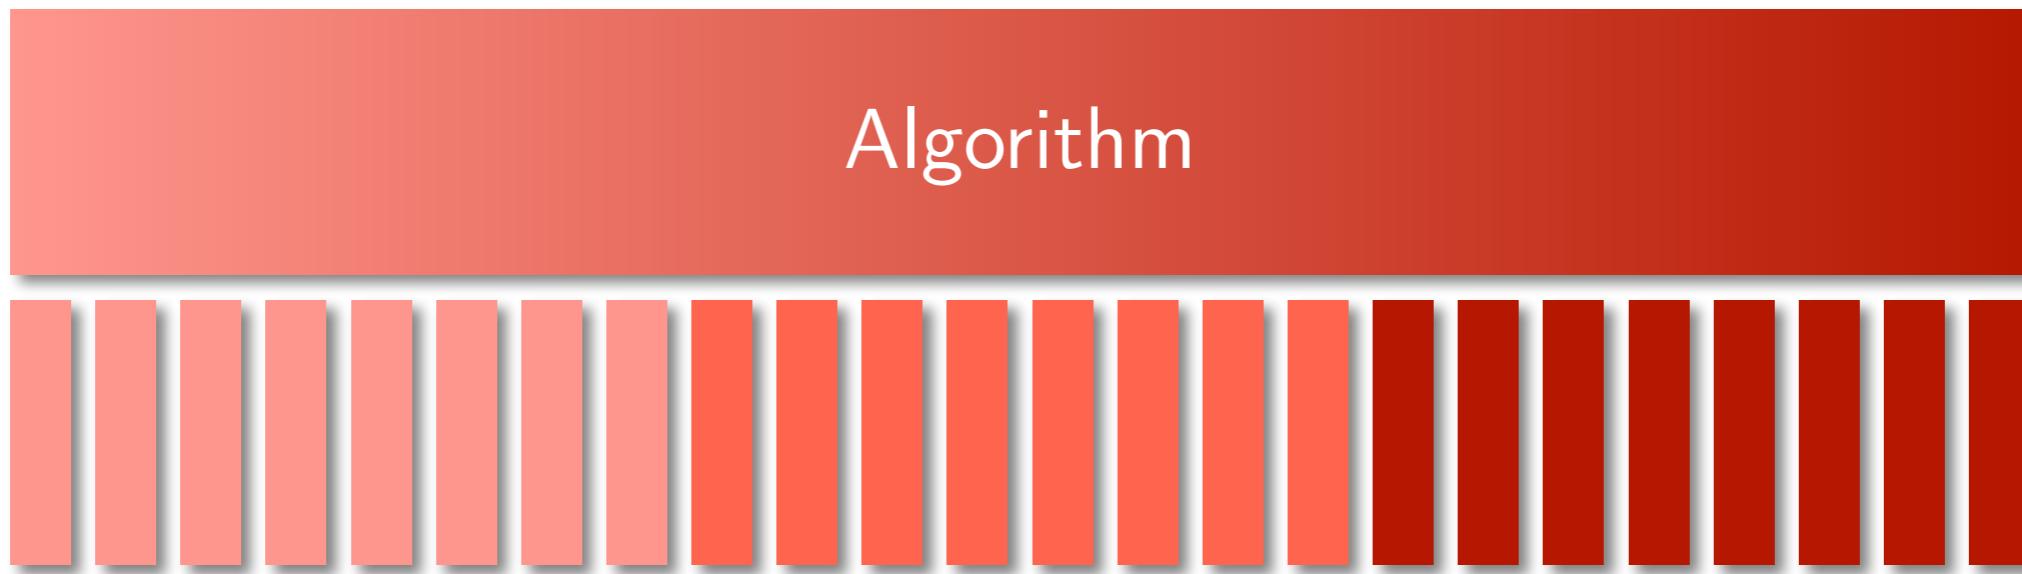


Data

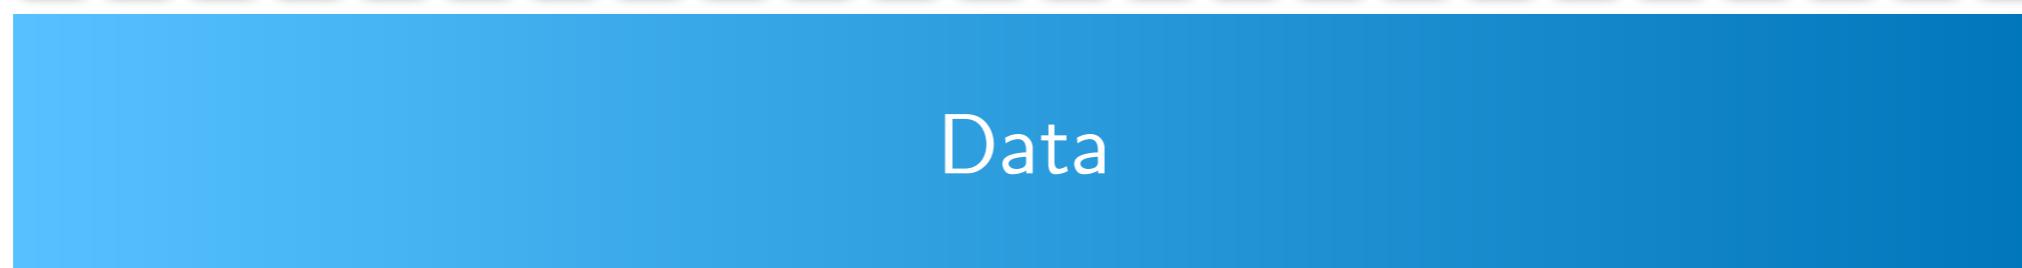
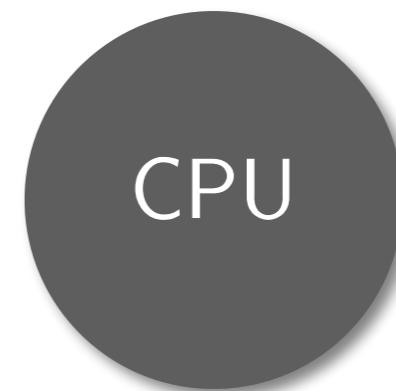
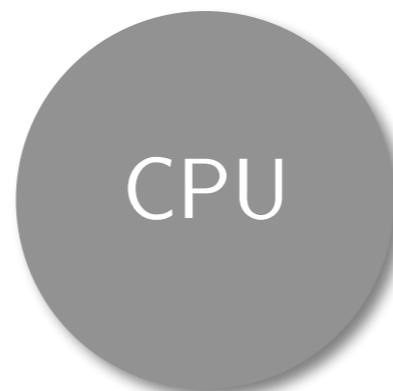
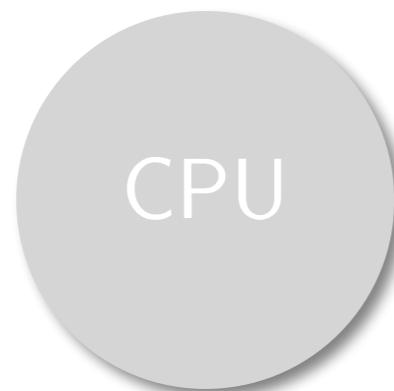
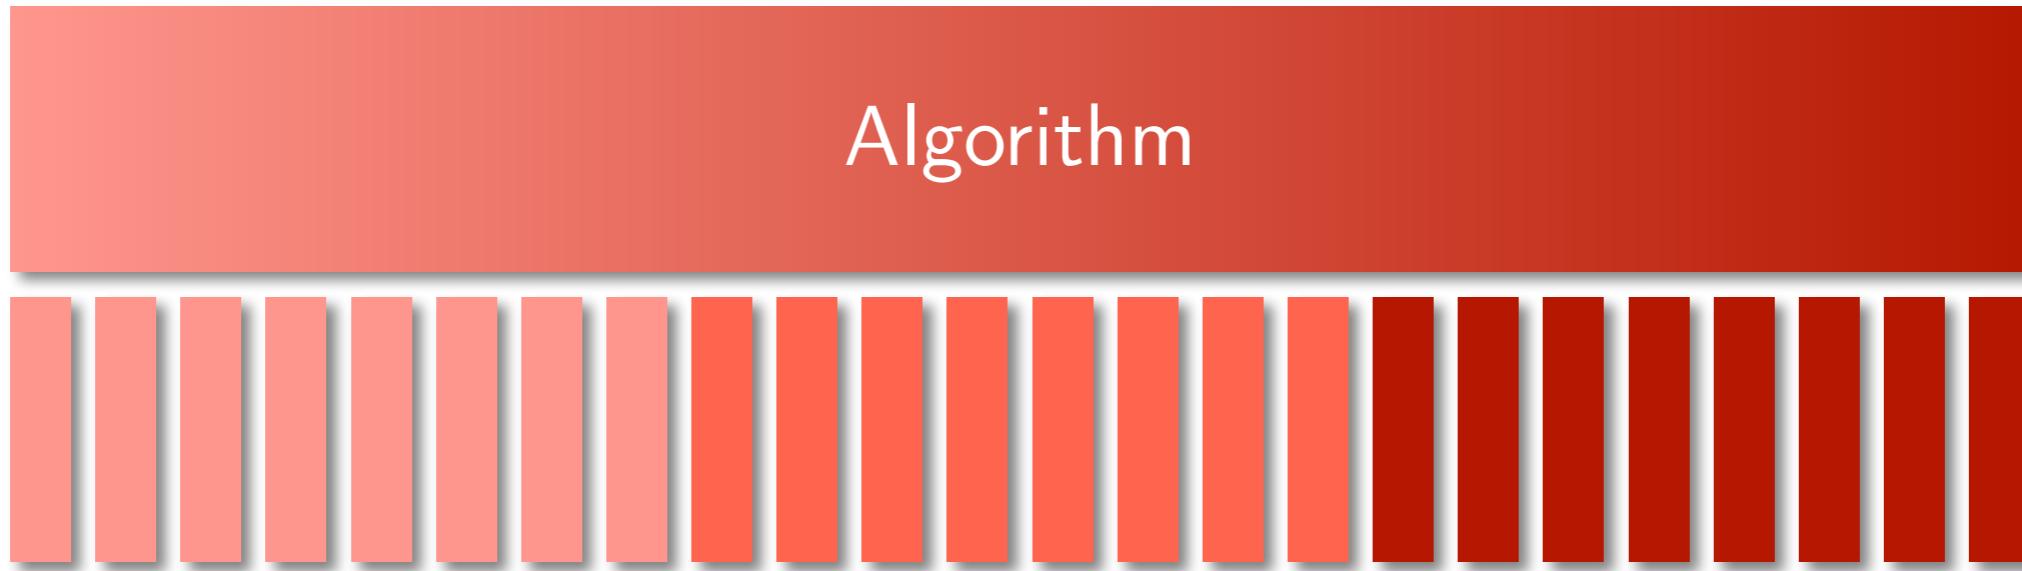
Single Instruction Multiple Data



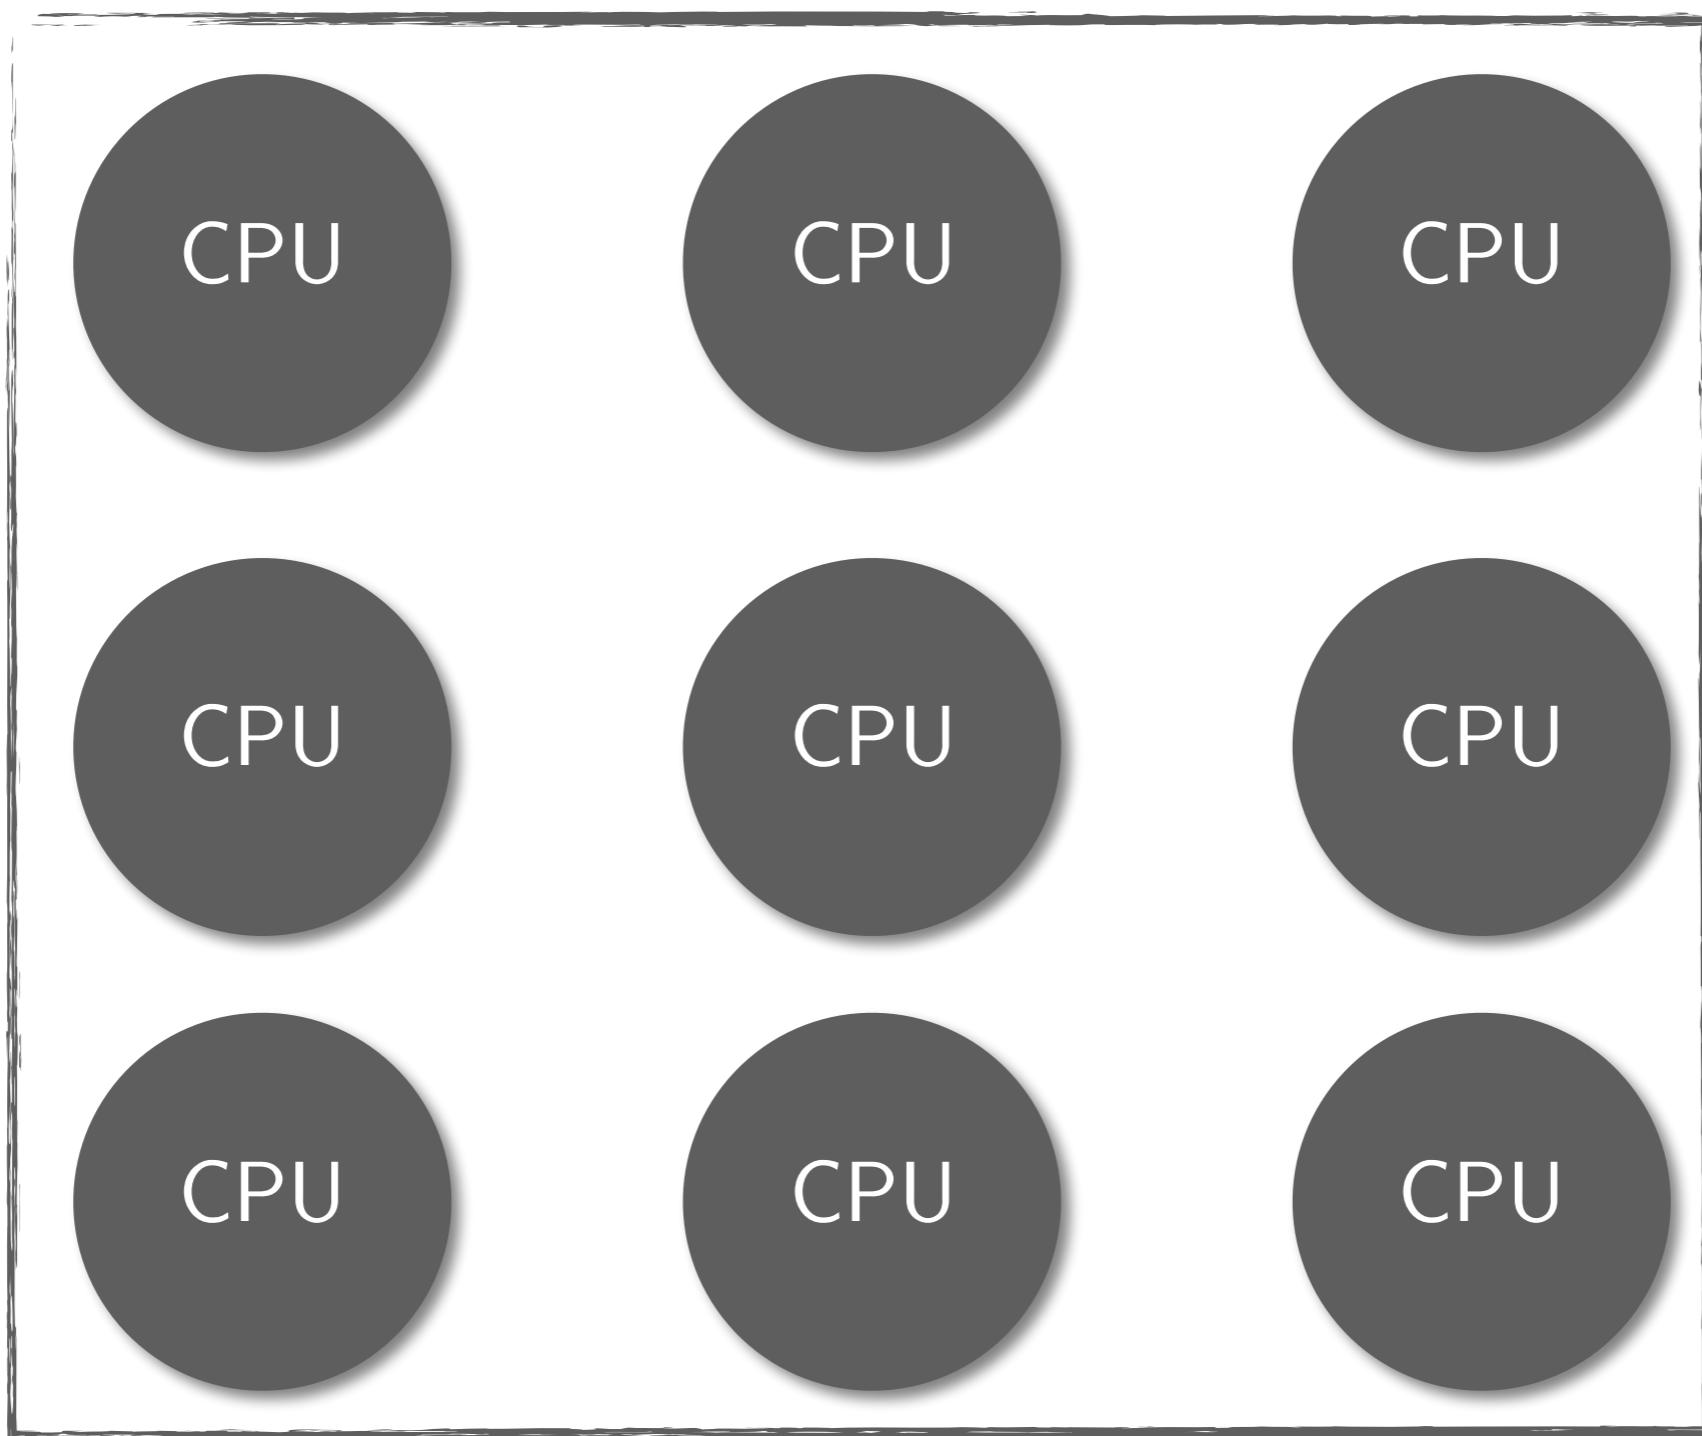
Multiple Instructions Single Data



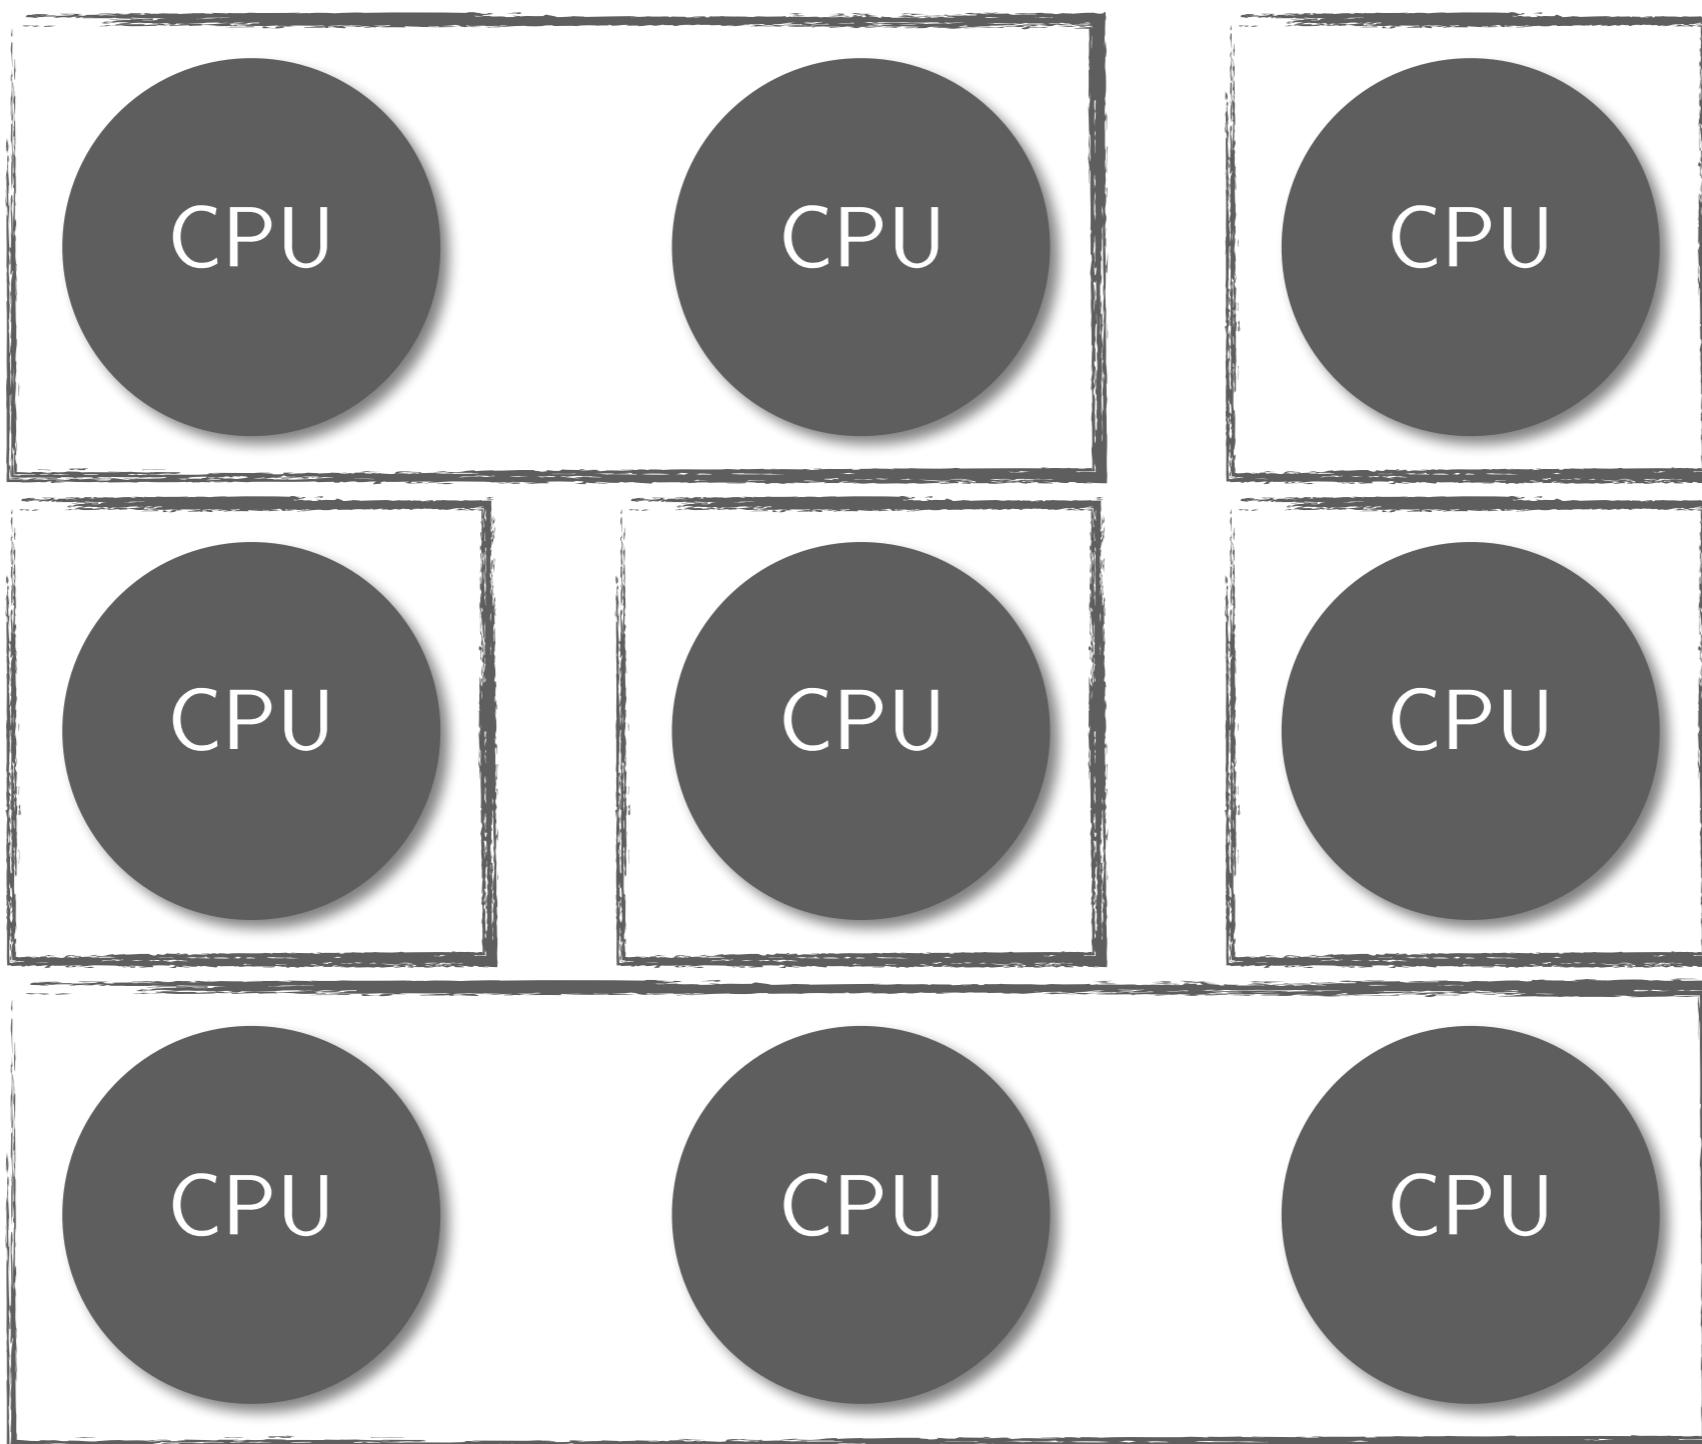
Multiple Instructions Multiple Data



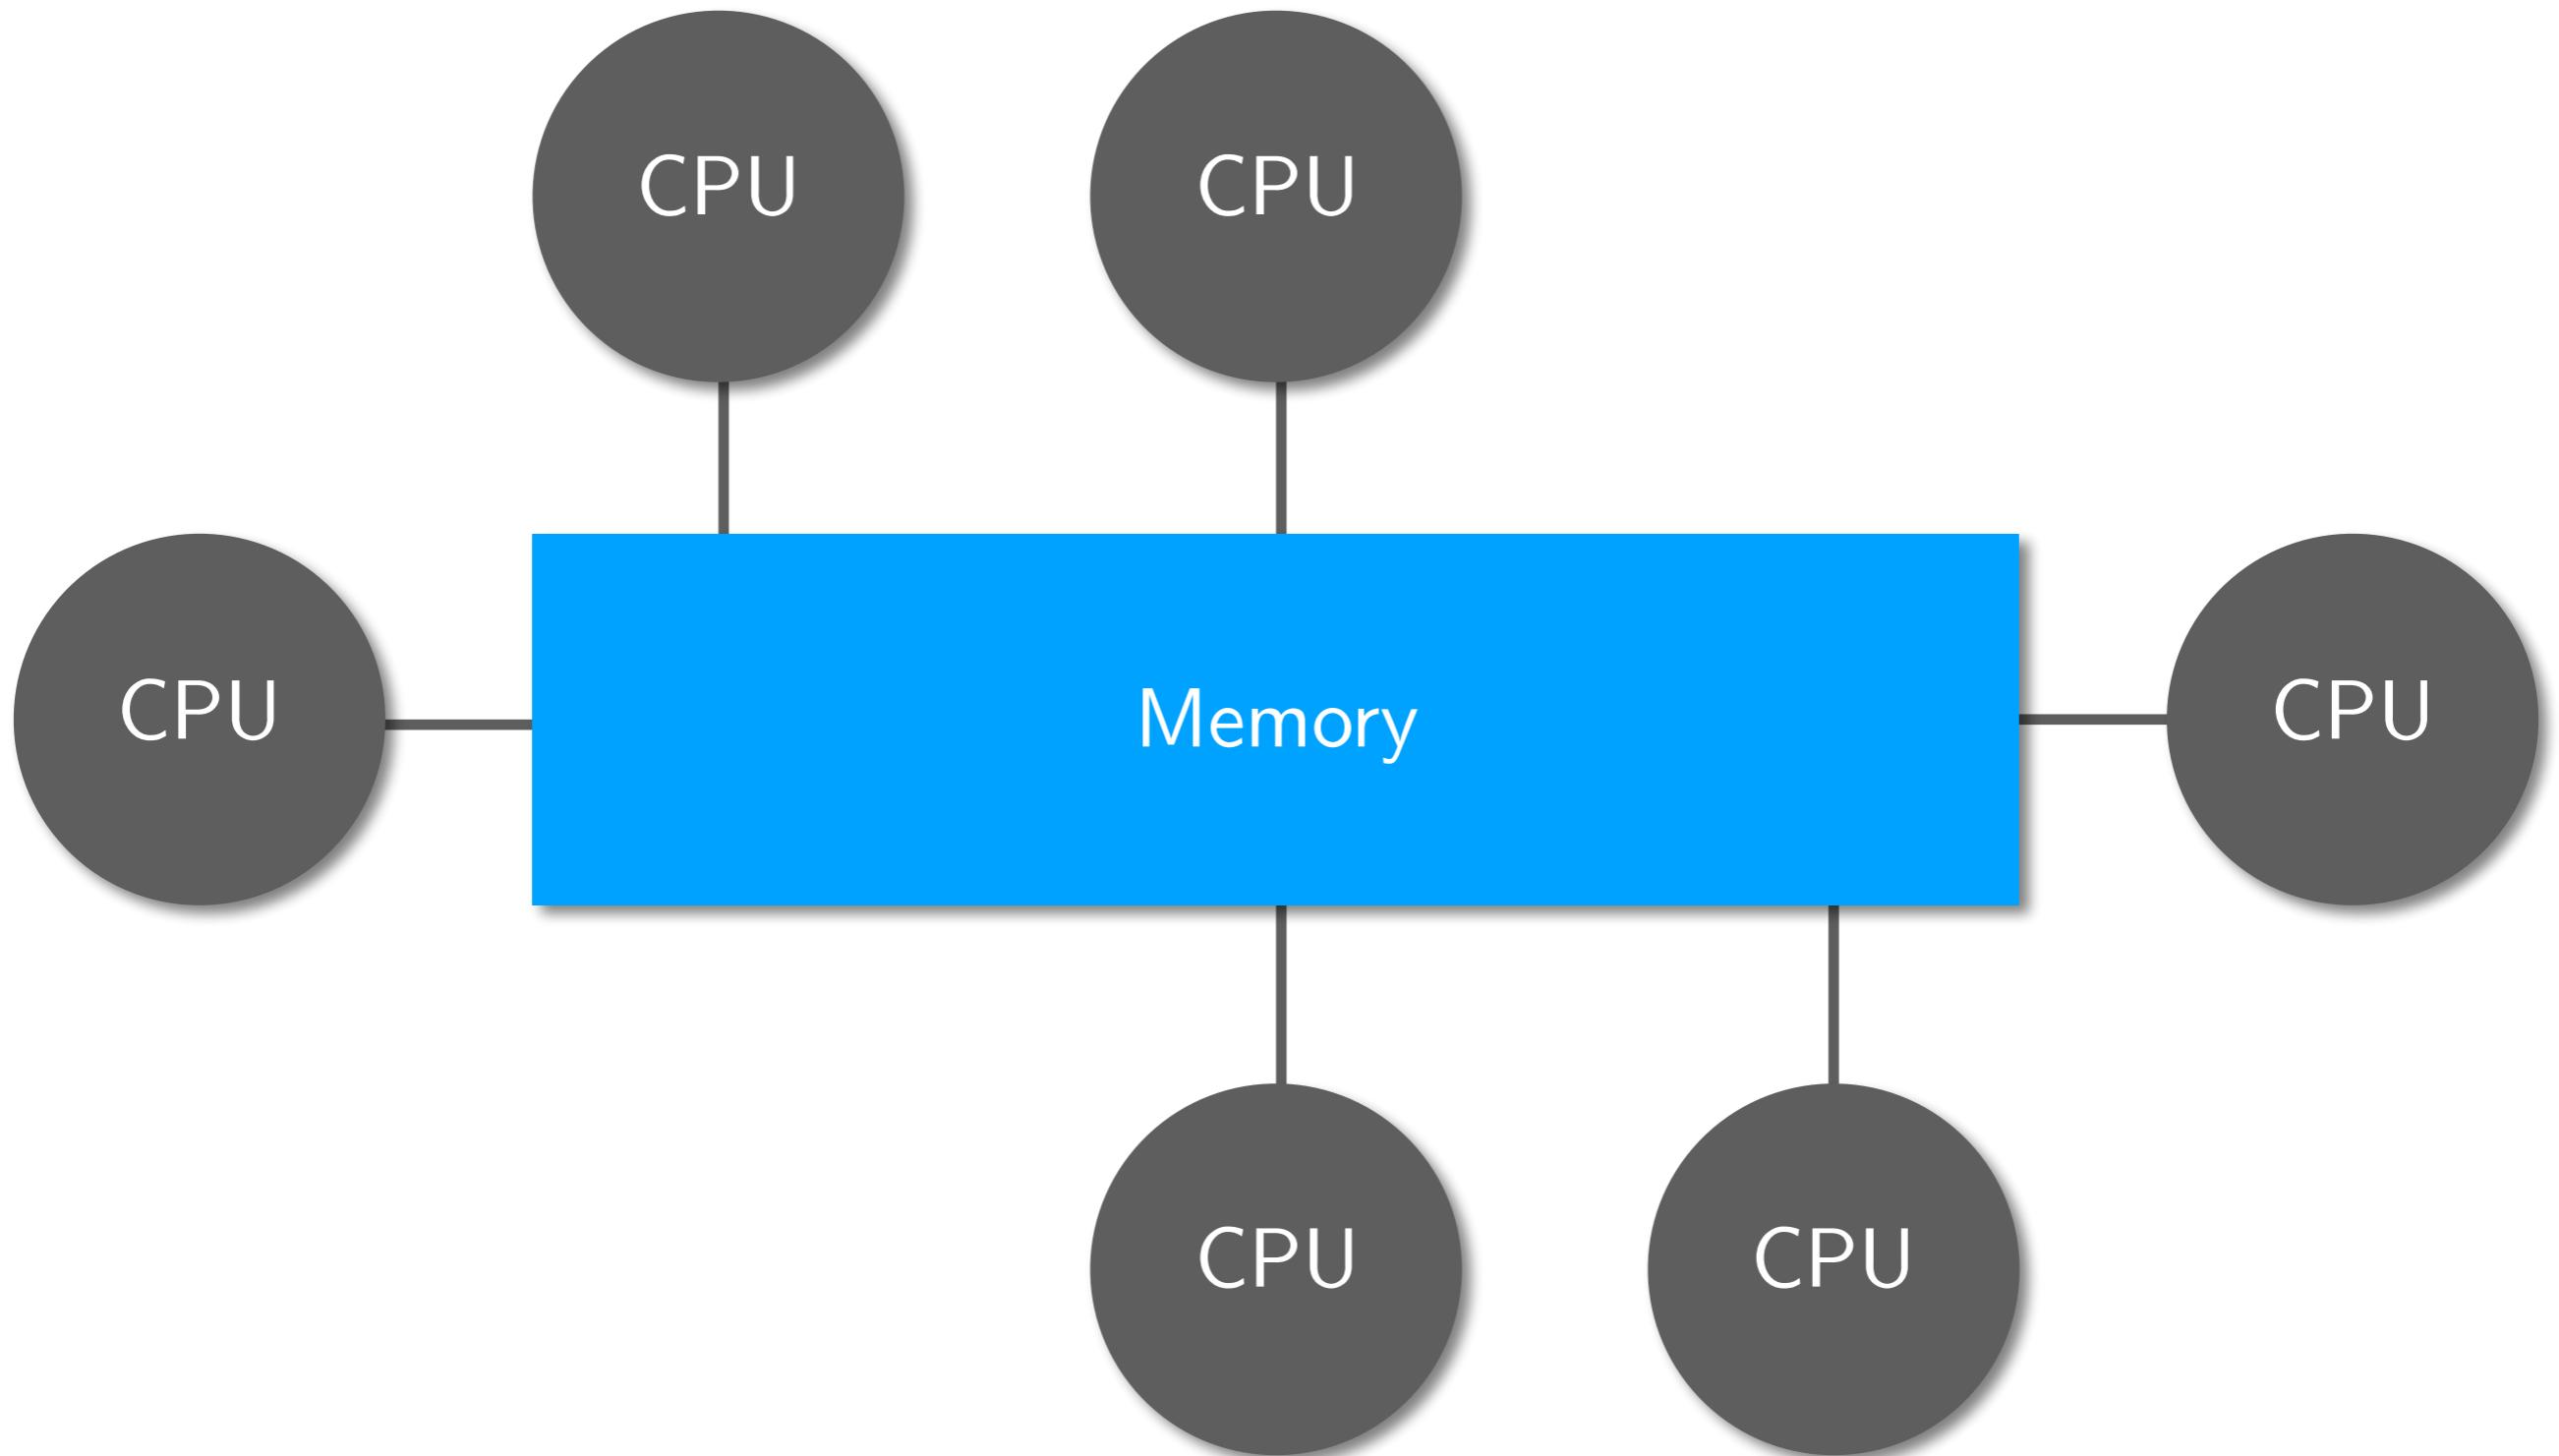
Parallel CPUs



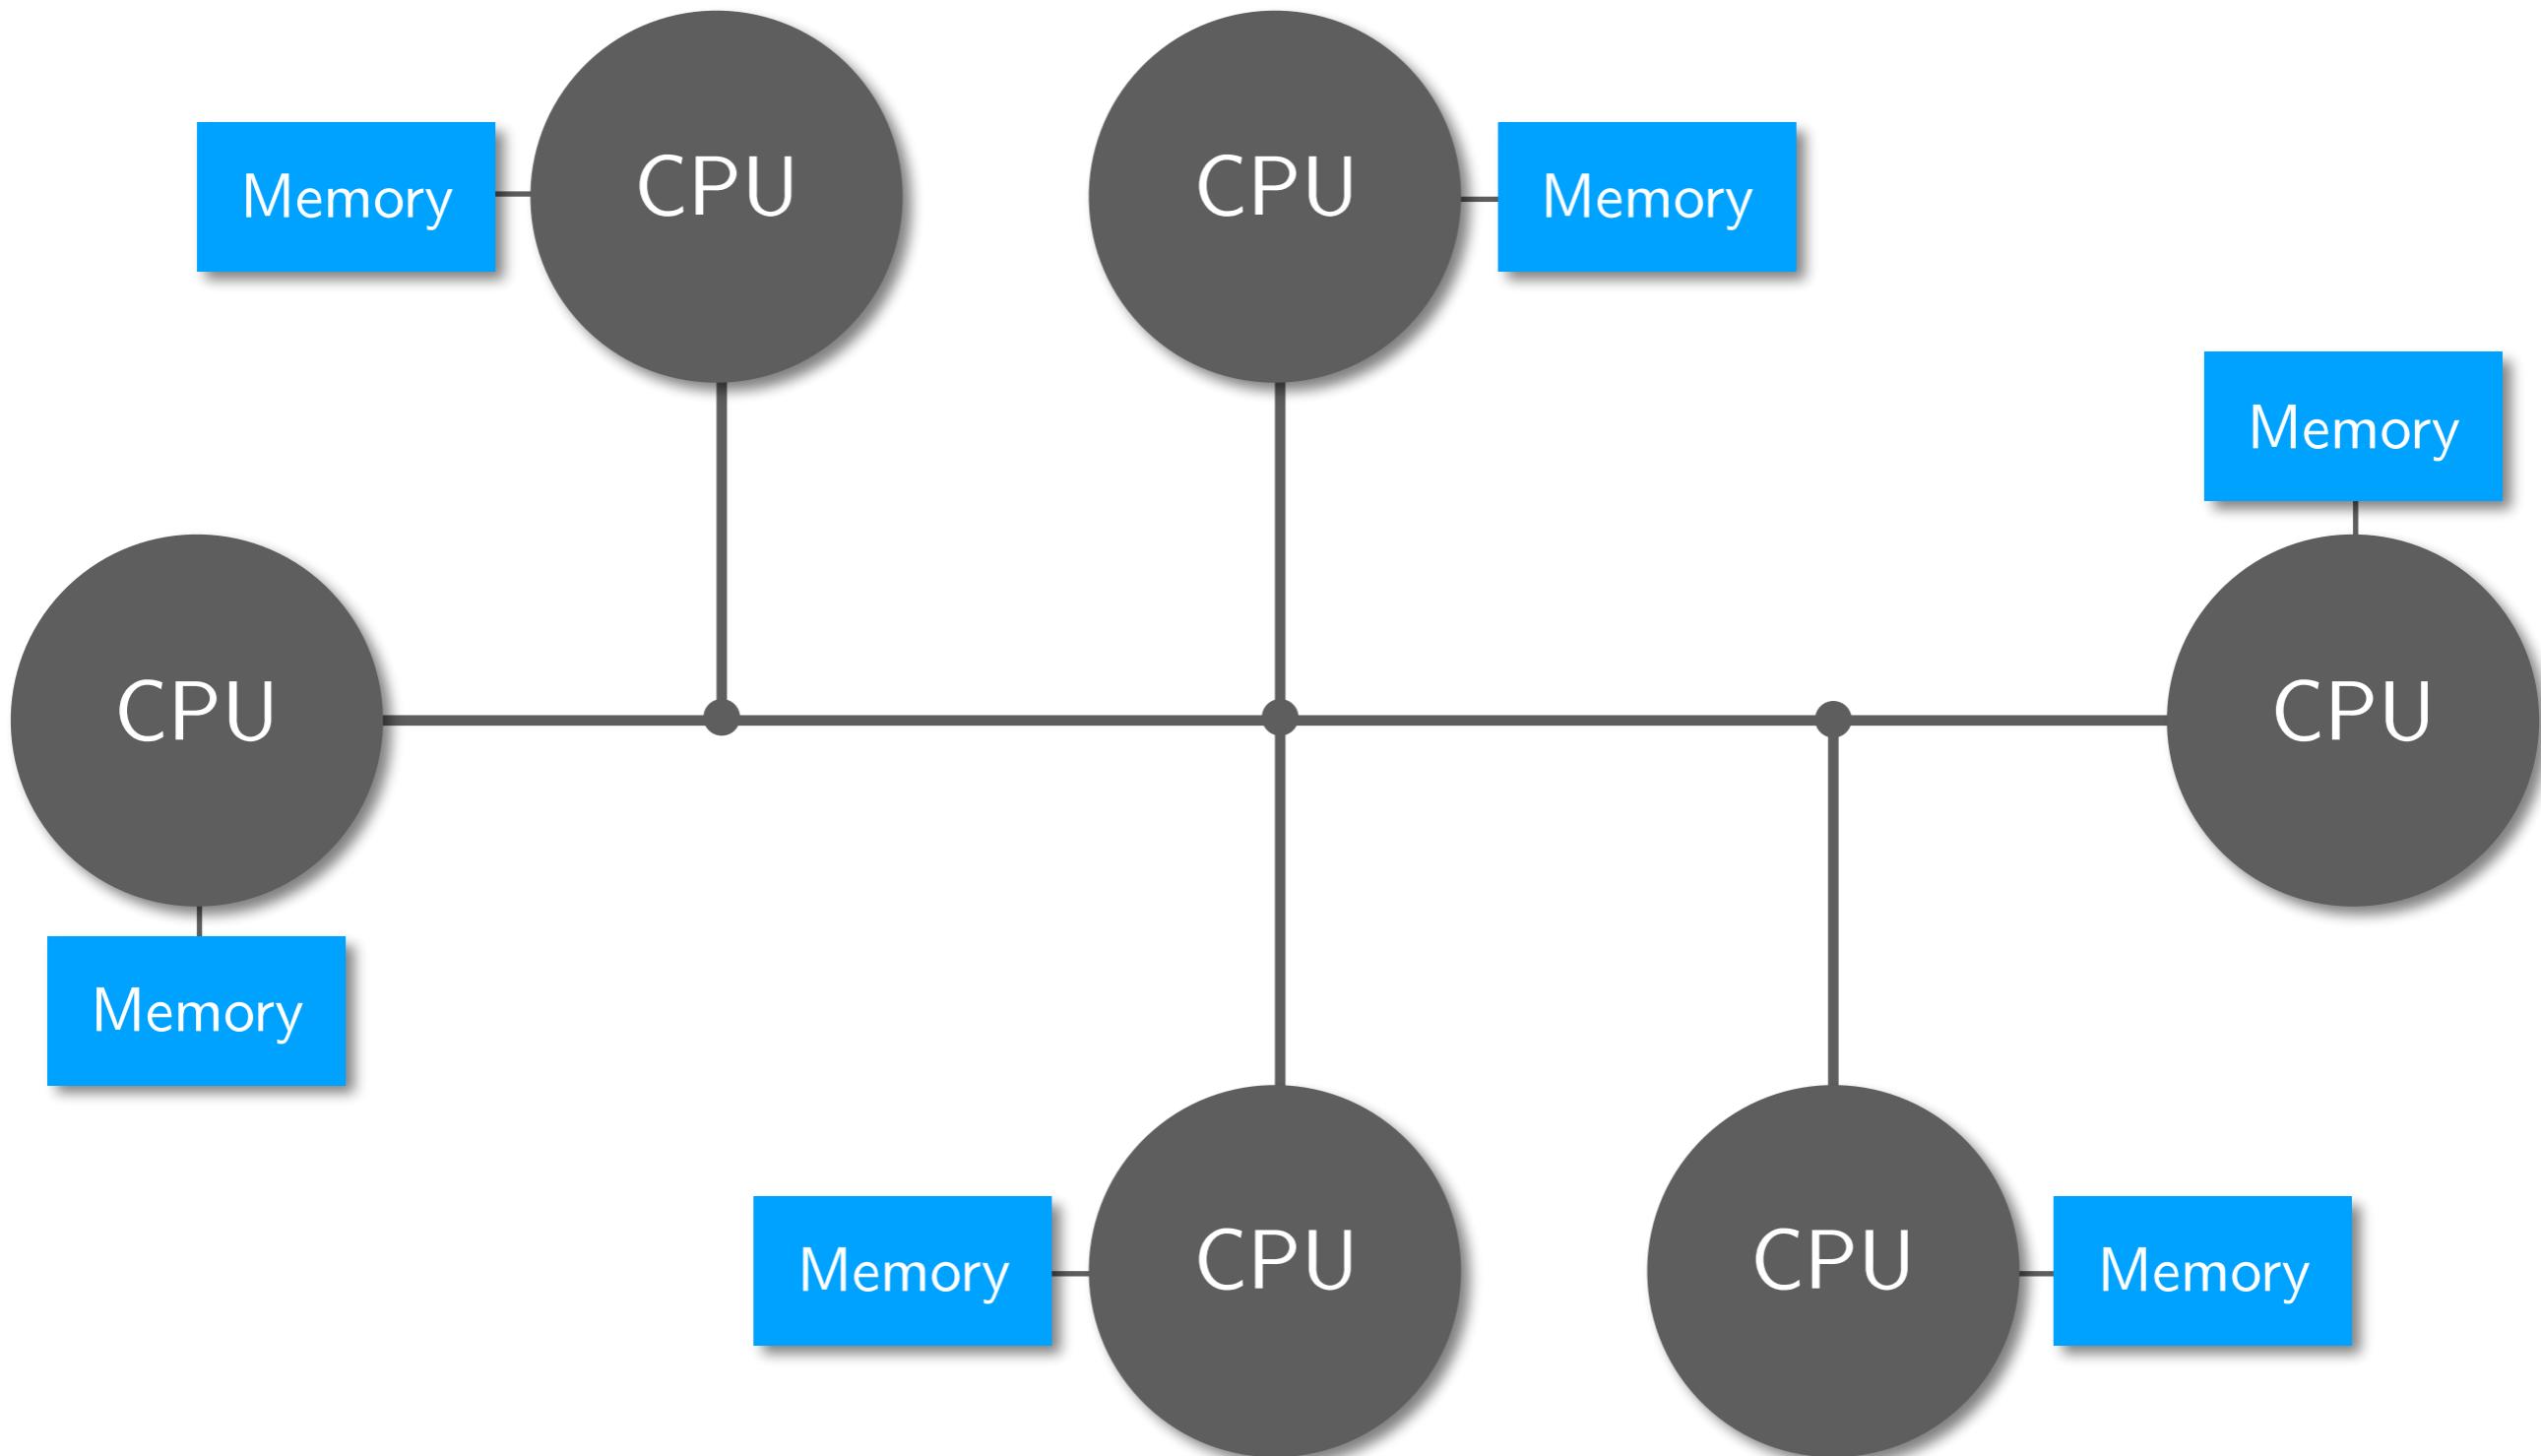
Distributed CPUs



Shared Memory



Distributed Memory



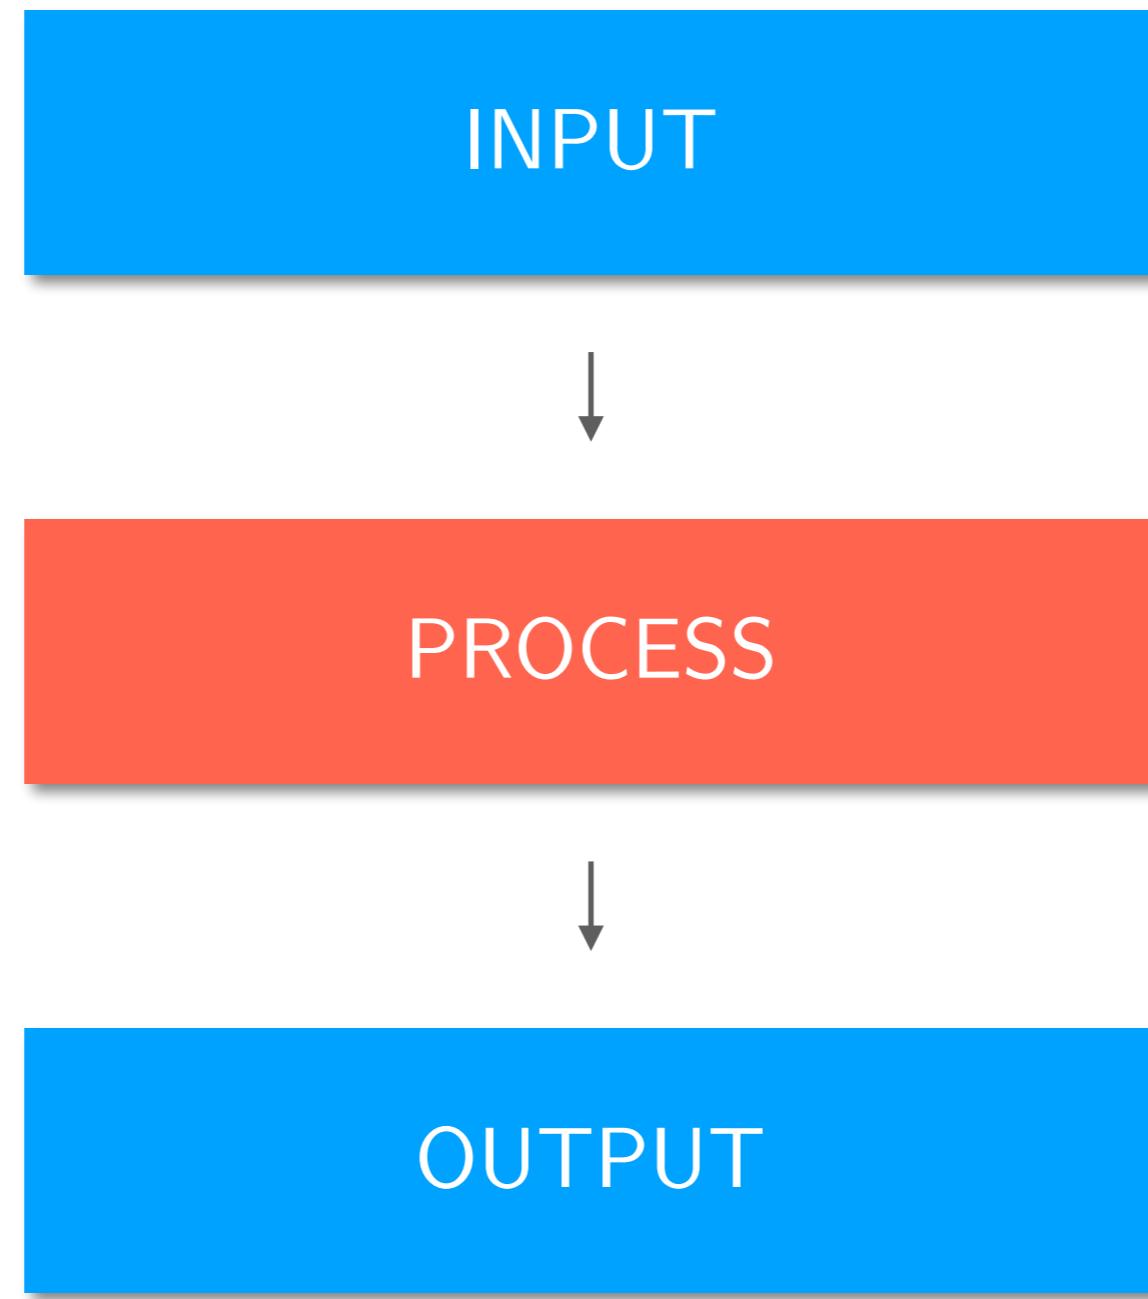
Splitting Code

- How to divide code into parallel tasks?
- How to distribute the code?
- How to coordinate the execution?
- How to load the data?
- How to store the data?
- What if more tasks than CPUs?
- What if a CPU crashes?
- What if a CPU is taking too long?
- What if the CPUs are different?
- What if we have a new (serial) code?

Splitting Data

- How to split the data?
- How to distribute the data?
- How to collect the data?
- How to merge the data?
- How to coordinate the access to the data?
- What if more data splits than tasks?
- What if tasks need to share data splits?
- What if a data split becomes unavailable?
- What if we have a new input?
- What if the data is big?

Typical Application



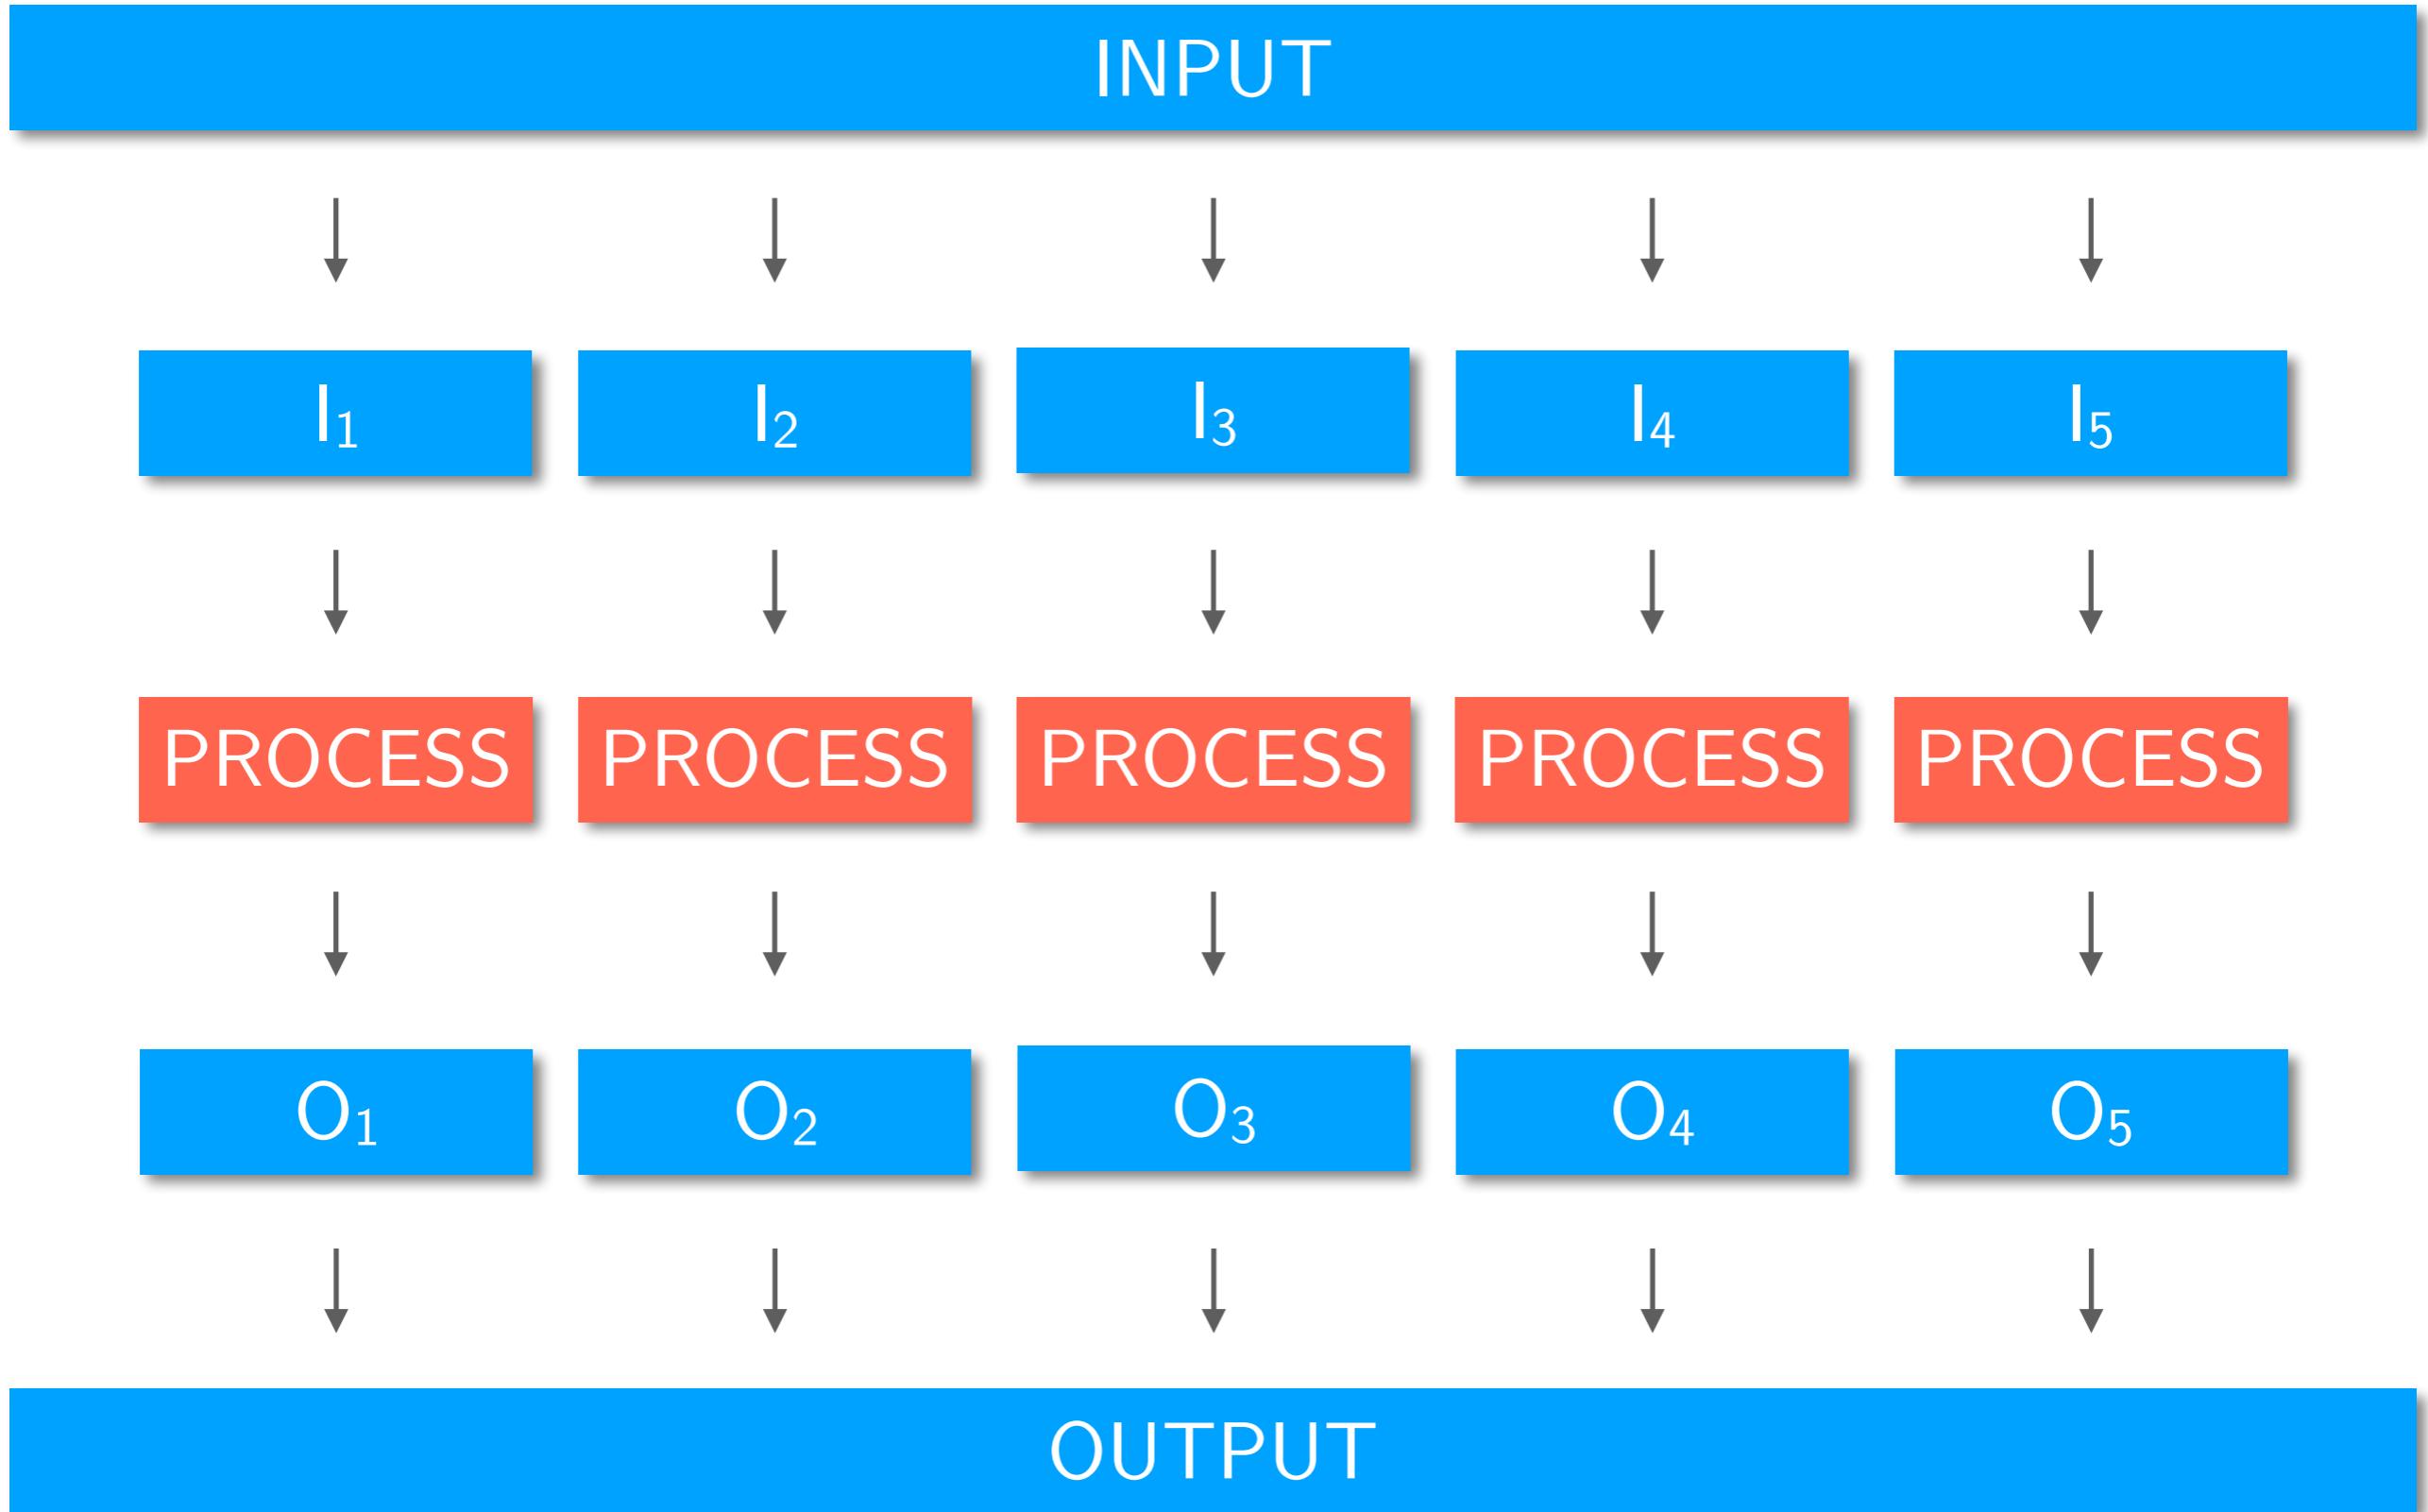
What if?

INPUT

PROCESS

OUTPUT

Divide and Conquer



Challenges

- How do we split the input?
- How do we distribute the input splits?
- How do we collect the output splits?
- How do we aggregate the output?
- How do we coordinate the work?
- What if input splits > num workers?
- What if workers need to share input/output splits?
- What if a worker dies?
- What if we have a new input?

Computing Platform



Source: <https://www.pexels.com/photo/datacenter-server-449401/>

Computing Platform



Source: <https://www.pexels.com/photo/datacenter-server-449401/>

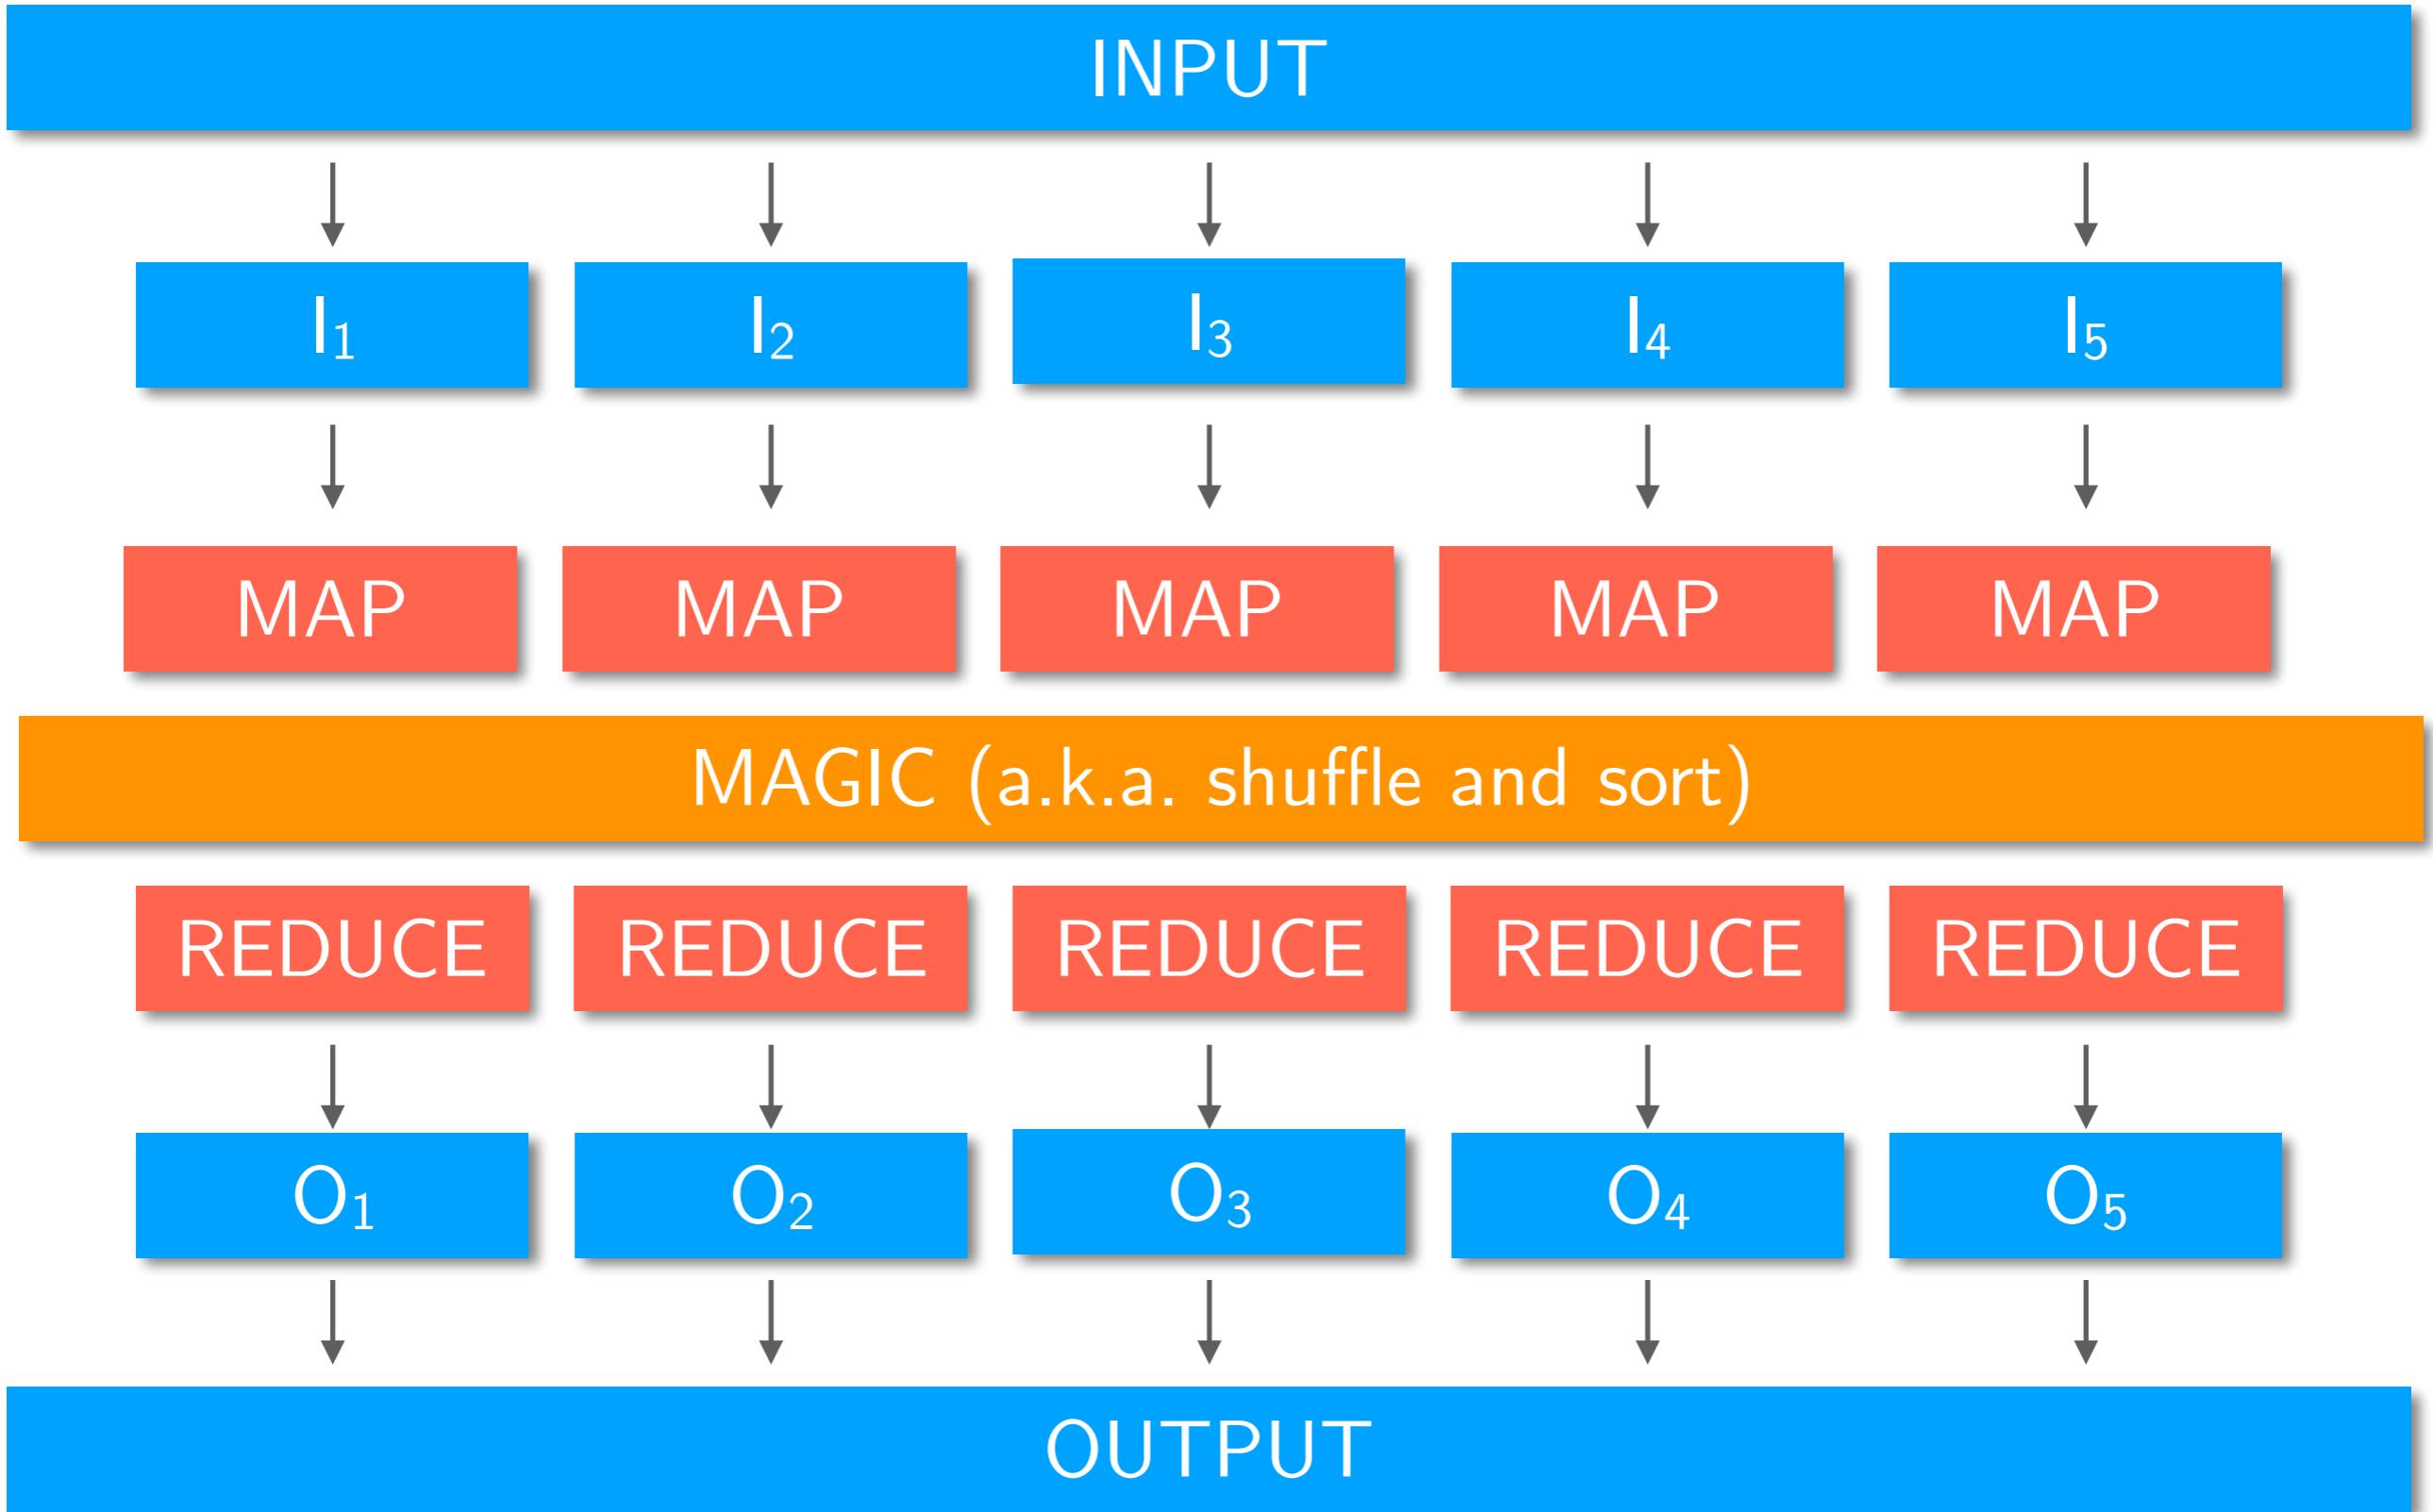
Design Ideas

- Scale “out”, not “up”
 - Avoid supercomputer, too costly
 - Use commodity machines, low costs
 - Drawback: many failures
- Move processing to the data
 - Same code, runs everywhere
 - Reduce data over the network
 - Drawback: code must be portable
- Process data sequentially, avoid random access
 - Huge data files (terabytes), not small files (megabytes)
 - Write once, read many
 - Drawback: poor support for standard file APIs
- Right level of abstraction
 - Hide implementation details from applications development
 - Write very few lines of code
 - Drawback: everything needs to fit into the abstraction

Typical Application

1. Iterate over a large number of records
2. Extract something of interest from each
3. Shuffle and sort intermediate results
4. Aggregate intermediate results
5. Generate final output

Map Reduce Application



Programming Model

- Programmers specify **two functions**
 - *map function*: from [key, value] (1) to [key, value] (0 or more)
 - *reduce function*: from [key, list of values] (1) to [key, values] (0 or more)
- **Map** function
 - Receives as input a key-value pair
 - Produces as output a list of key-value pairs (typically 1 or more per input)
- **Reduce** function
 - Receives as input a key-list of values pair
 - Produces as output a list of key-value pairs (typically none or 1 per input)
- Both functions are **STATELESS**
- The **runtime support** handles everything else...

Wordcount Example (I)

```
class MAPPER
    method MAP(docid a, doc d)
        for all term t in doc d do
            EMIT(term t, count 1)

class REDUCER
    method REDUCE(term t, counts [c1, c2, ...])
        sum ← 0
        for all count c in counts [c1, c2, ...] do
            sum ← sum + c
        EMIT(term t, count sum)
```

Wordcount Example

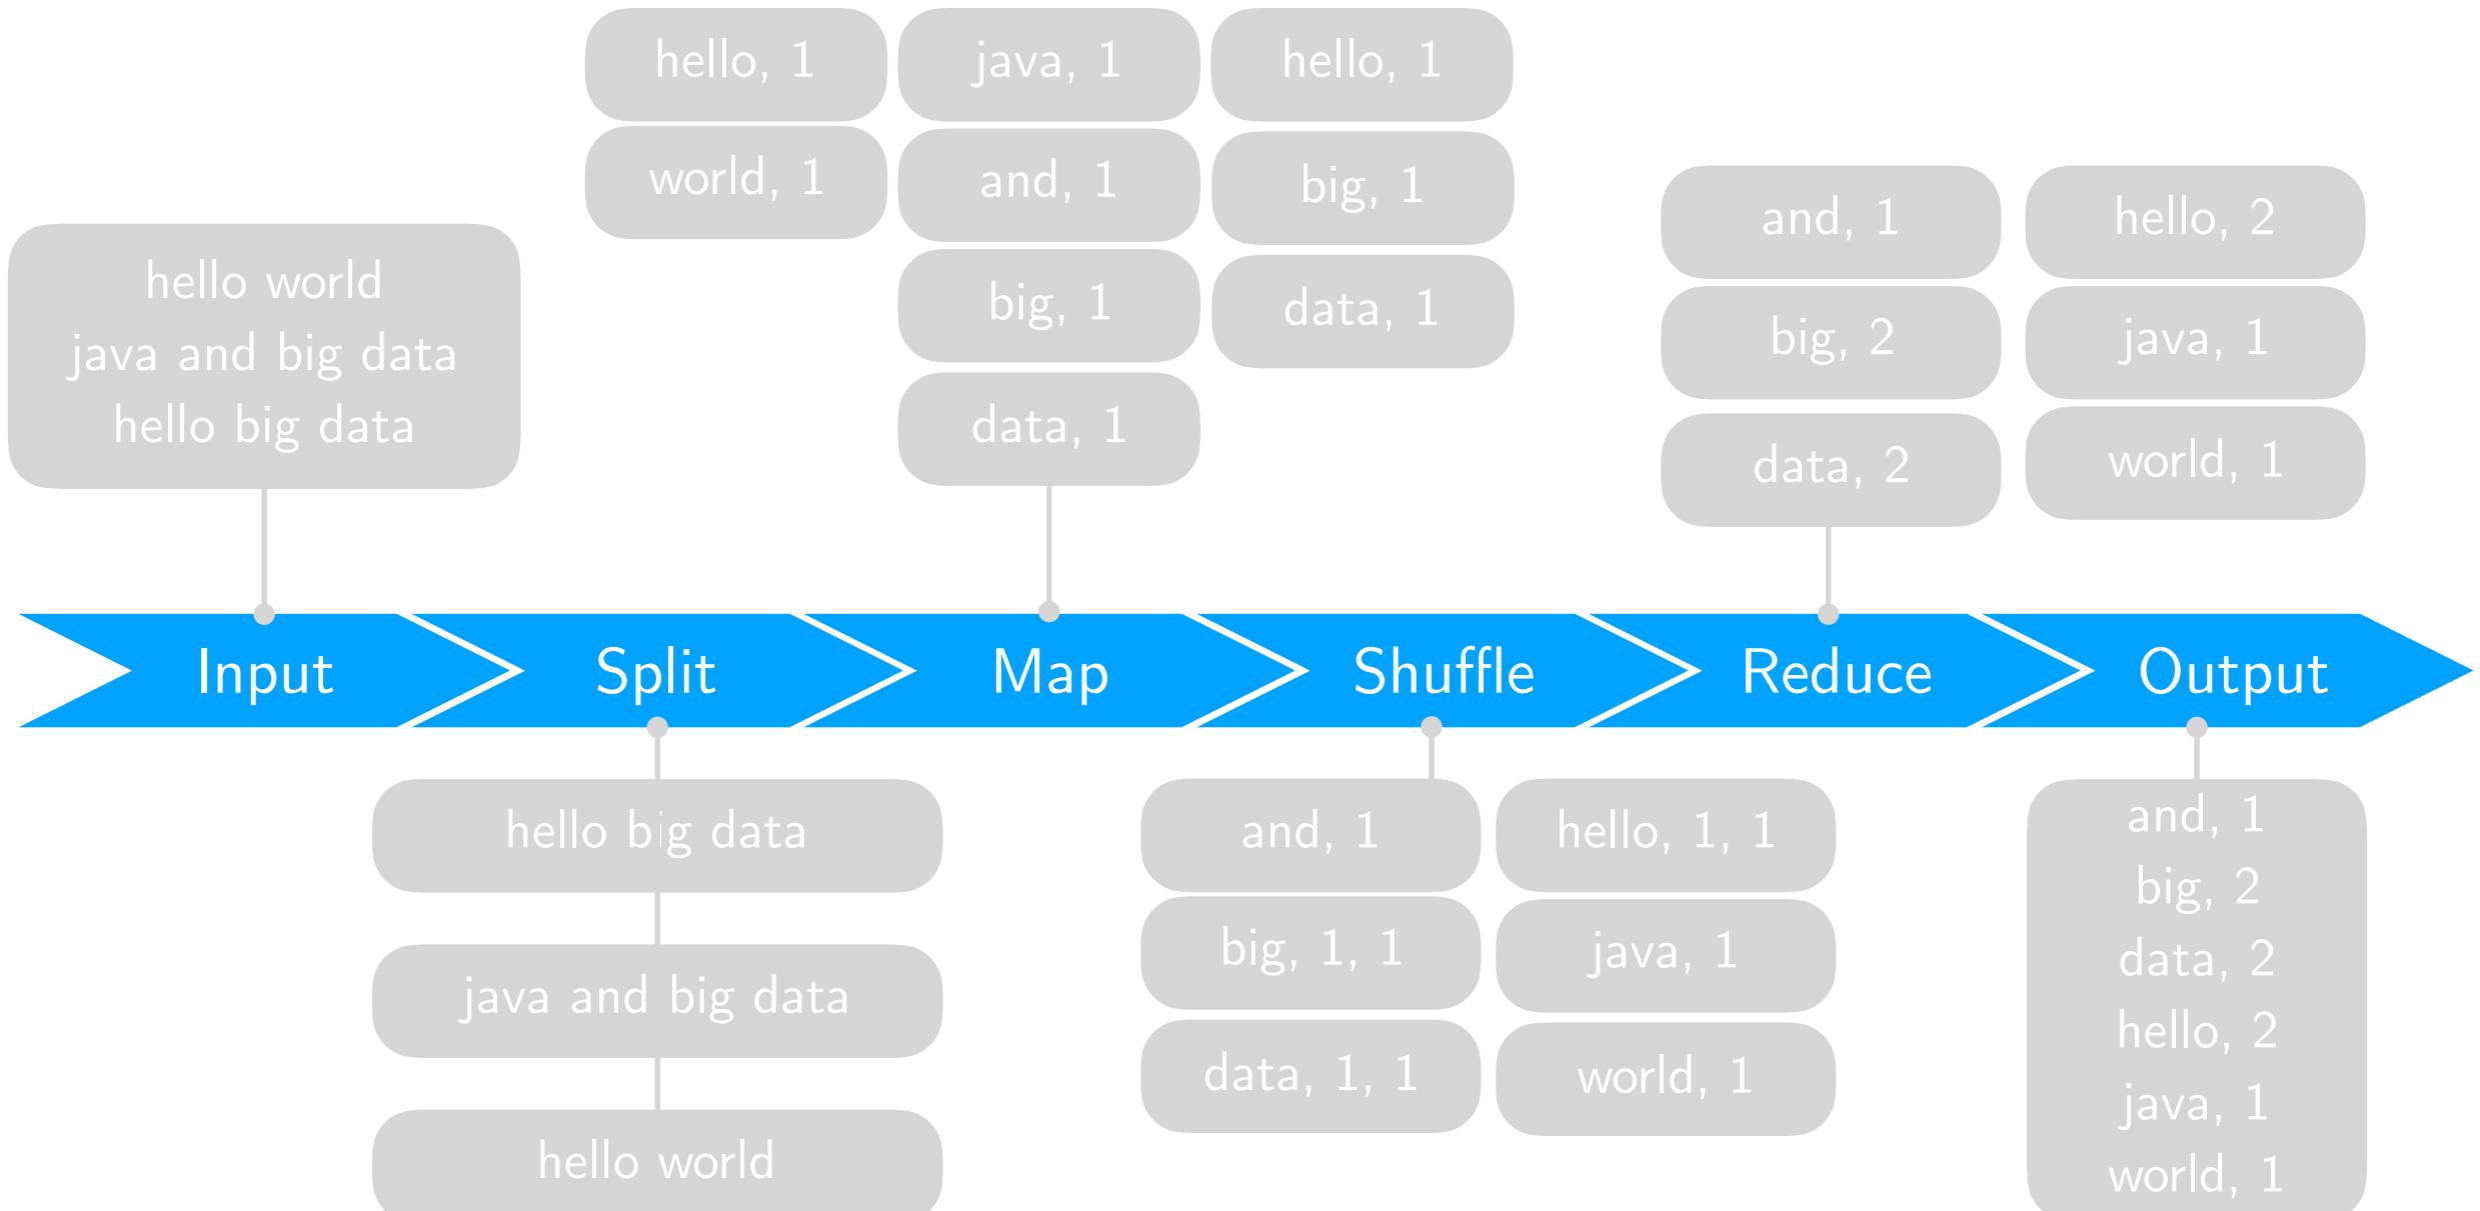


Image Data Example (I)

- Image data from **different** content providers
 - Different formats
 - Different coverages
 - Different timestamps
 - Different resolutions
 - Different exposures/tones
- **Large amount** to data to be processed
- Goal: produce data to serve a "**satellite**" view to users

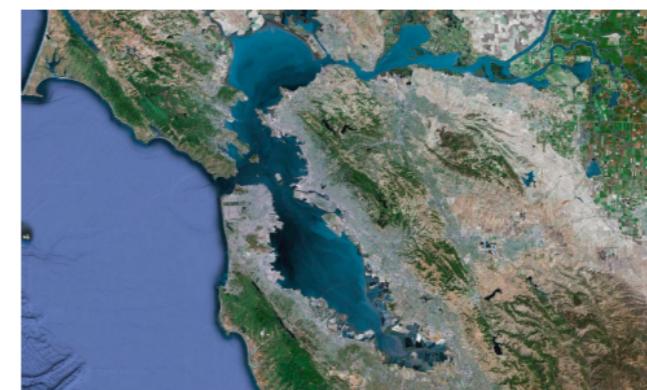
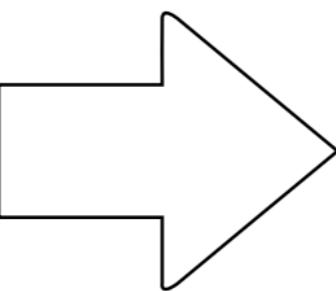
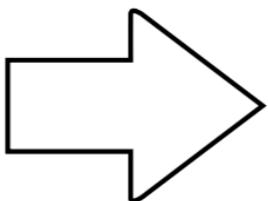
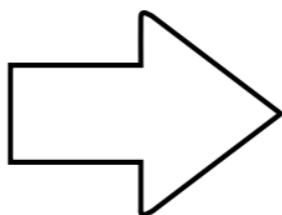


Image Data Example (II)

- Split the whole territory into "tiles" with fixed location IDs
- Split each source image according to the tiles it covers



- For a given tile, stitch contributions from different sources, based on its freshness and resolution, or other preference

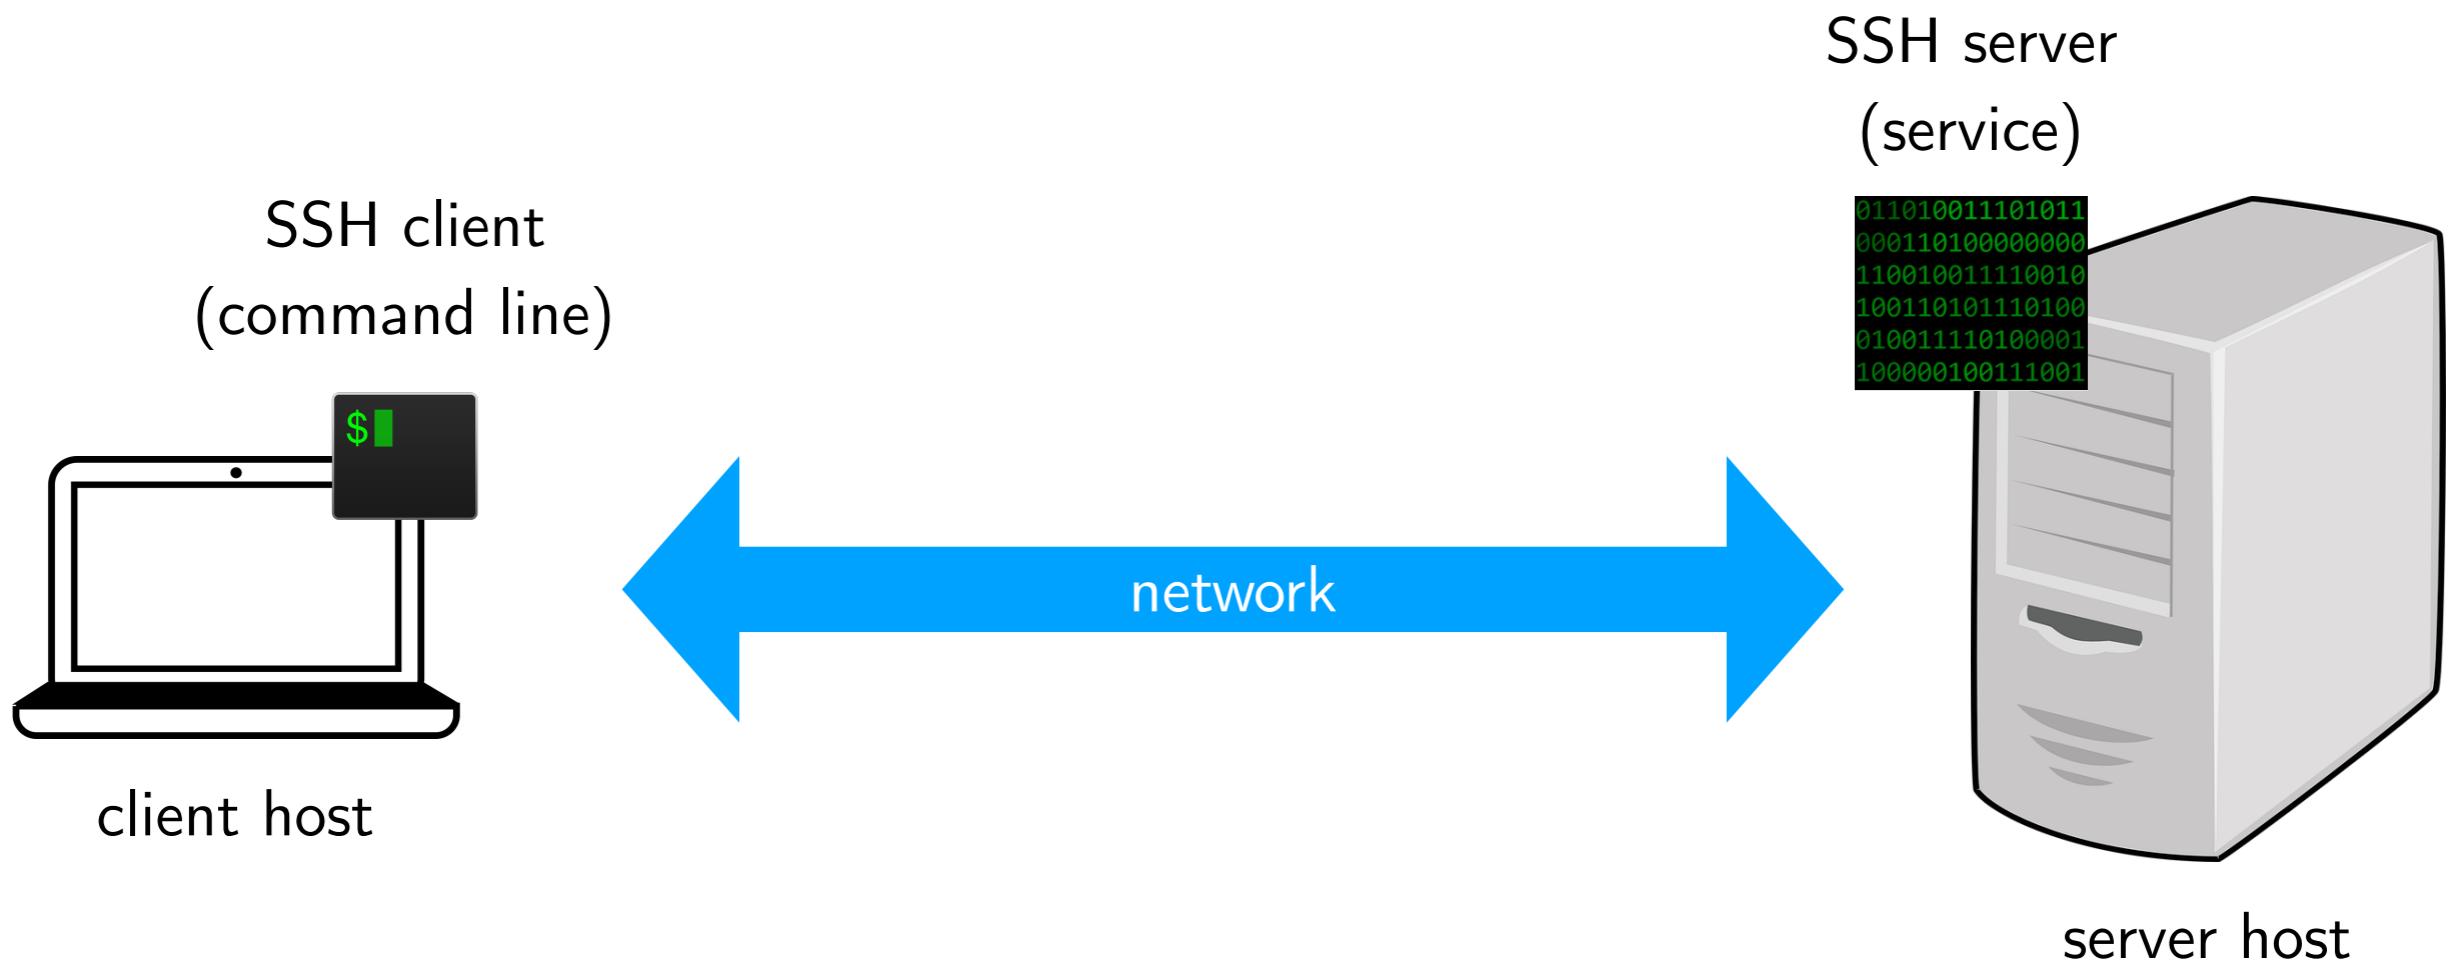


- Serve the merged imagery data for each tile, so they can be loaded into and served from an image server farm.

SSH Crash Course

- SSH = Secure Shell
- Replacement for RSH, the Unix remote shell application
 - RSH allowed one to connect to a shell on a remote machine
 - First issue: all traffic in cleartext
 - Second issue: access to a machine grants access to all machines
 - Command-line connection tool
 - All traffic encrypted
 - Both ends authenticate themselves to the other end

SSH Picture



SSH Basic Use

- ssh <server_name>
- ssh -l <user_name> <server_name>
- ssh <user_name>@<server_name>
- ssh <server_name> <command>
- ssh -v <server_name>

Secure File Copy

- SCP = Secure Copy

```
scp localfile remotemachine:/remotepath/file
```

- The local and remote file specifications are of the form

```
user@host:/path/to/file
```

- If omitted, user, host and path/to/file directory default to local username, local hostname, and current directory, respectively.
- Prompts for authentication if needed
- All traffic encrypted
- If including wildcards (*), include the block in quotes ("")